

Project Report: Course Project-3

CS51510-001, Fall 2025
Purdue University Northwest
12/12/2025

Class ID (Ascend sequency)	Contributor Name (Last Name, First Name)	Email	Detailed Contributions
038414971	Mishra, Amartya	Mishr259@pnw.edu	Documentation
039133149	Mammai, Sreeja	Smmamai@pnw.edu	Data Collection
038410044	Mhasawade, Siddhi	Smhasawa@pnw.edu	Performance Testing
0038159254	Mekala, Ruthvik	Rmekala@pnw.edu	Python Code and UI
039133149	Patel, Raaj	Pate2682@pnw.edu	Research Report and Architecture Diagram
0037874079	Modi, Hirak	Modi54@pnw.edu	Project Presentation
0038119426	Nalluri, Prasanth	Pnallur@pnw.edu	Python Code and UML Diagrams

Table of Contents

Abstract	4
I. Introduction	6
II. Methodology	9
1. Data Collection	9
i. City Attributes	9
ii. Inter-City Map Distances.....	10
iii. Weather-Risk Conditions	12
iv. Operational Constraints and Fuel Model	14
2. System Design	15
i. Backend Computation Engine	15
ii. API Layer and Routing Interface.....	16
iii. Frontend Visualization Engine	17
iv. Integration Strategy.....	17
3. Graph Algorithms Implemented	19
i. Breadth-First Search (BFS)	19
ii. Depth-First Search (DFS)	20
iii. Prim's Minimum Spanning Tree Algorithm	20
iv. Kruskal's Minimum Spanning Tree Algorithm.....	21
v. Bellman–Ford Shortest-Path Algorithm	21
4. UML and ER Diagram Representation.....	23
5. Travel Computation Models	25
i. Elevation-Adjusted Distance Model	25
ii. Gasoline Consumption Model	26
iii. Weather-Risk Scoring Model	26
iv. Daily Driving Constraints and Multi-Day Risk Evaluation.....	27
v. Best Travel Date Identification	27
vi. Combined Route Scoring Function.....	28
III. Experimental Results	29
IV. Results and Analysis.....	40
1. Performance Monitoring on Linux	40
V. Reproducibility and code accessibility	51
VI. Conclusion	53
Acknowledgement	54

References	55
Appendix	57
1. Python Program	58

Table of Figures

Figure 1 - System Overview Diagram	7
Figure 2 — Attributes of 28 cities	9
Figure 3 — Map distance from source to destination.....	11
Figure 4 - Monthly data of weather risk for Chicago	13
Figure 5 - The entire pipeline from taking inputs to visualization of Data.....	18
Figure 6 - City Selection Drop-Down Interface	30
Figure 7 - Algorithm Selection UI.....	31
Figure 8 - Best Route Visualization.....	32
Figure 9 - BFS Route Visualization.....	33
Figure 10 - DFS Route Visualization.....	34
Figure 11 - Prim MST Route Visualization.....	35
Figure 12 - Kruskal MST Route Visualization.....	36
Figure 13 - Bellman–Ford Route Visualization.....	37
Figure 14: Python code running on ubuntu terminal	41
Figure 15: Fetching Process ID.....	42
Figure 16: Inspecting CPU Usage.....	43
Figure 17: CPU Percentage.....	44
Figure 18: free -m	44
Figure 19: vmstat command.....	45
Figure 20: pmap to identify memory segments	46
Figure 21	46
Figure 22: Output for iostat command.....	47
Figure 23	47
Figure 24: pidstat	48
Figure 25: uptime	48
Figure 26: process tree and state	49

Abstract

This project develops an integrated travel-route optimization framework that combines classical graph algorithms with real-world environmental, geographic, and operational constraints. The system constructs a weighted, elevation-aware transportation graph across 28 U.S. cities, where each edge incorporates Google-Maps derived distances, sea-level differences, and a physics-based adjustment using $\tan(\theta)$ to represent fuel penalties associated with incline. To reflect realistic travel conditions, the model incorporates daily weather-risk values for each city during November 1-30, 2025, and computes the day with the minimum accumulated risk along any proposed route. The platform further estimates fuel consumption, enforces speed and daily-driving limitations, and derives a combined route-quality score that balances distance, fuel usage, and safety risk.

Five fundamental algorithms Breadth-First Search (BFS), Depth-First Search (DFS), Prim's Minimum Spanning Tree, Kruskal's Minimum Spanning Tree, and the Bellman–Ford shortest-path algorithm were implemented in Python and exposed through a lightweight HTTP API. A browser-based visualization layer built with D3.js renders the U.S. map, city nodes, and route segments, allowing users to compare algorithmic behaviours interactively. Extensive experimentation highlights the differing structural assumptions of each algorithm, their sensitivity to edge-weight models, and their impact on total distance, gas consumption, and weather-risk

exposure. Performance profiling on a Linux server further evaluates CPU, memory, and I/O behaviour during large-scale route computations.

Beyond demonstrating algorithmic correctness, this project emphasizes the importance of integrating contextual data into route planning. Real-world travel is rarely optimized by distance alone; elevation, fuel cost, and regional weather patterns may dramatically alter what constitutes a “best” route. By combining classical graph traversal with empirical risk metrics and environmental modelling, the system illustrates how computational methods can be expanded to deliver more realistic and safety-conscious travel recommendations. This approach offers a blueprint for broader applications in logistics, disaster planning, and intelligent transportation systems.

I. Introduction

Route planning is a fundamental problem in transportation science, logistics, and computer science. Classical graph algorithms provide an elegant theoretical foundation for computing paths, yet real-world travel rarely aligns with the simplified conditions assumed in textbook formulations. Distances fluctuate with terrain, fuel consumption depends on elevation changes, and environmental factors such as rain, snow, and cloud cover may influence the safety and practicality of travel on any given day. The central motivation behind this project is to bridge this gap between theoretical path computation and realistic travel decision-making by integrating heterogeneous, real-world datasets into a unified algorithmic framework.

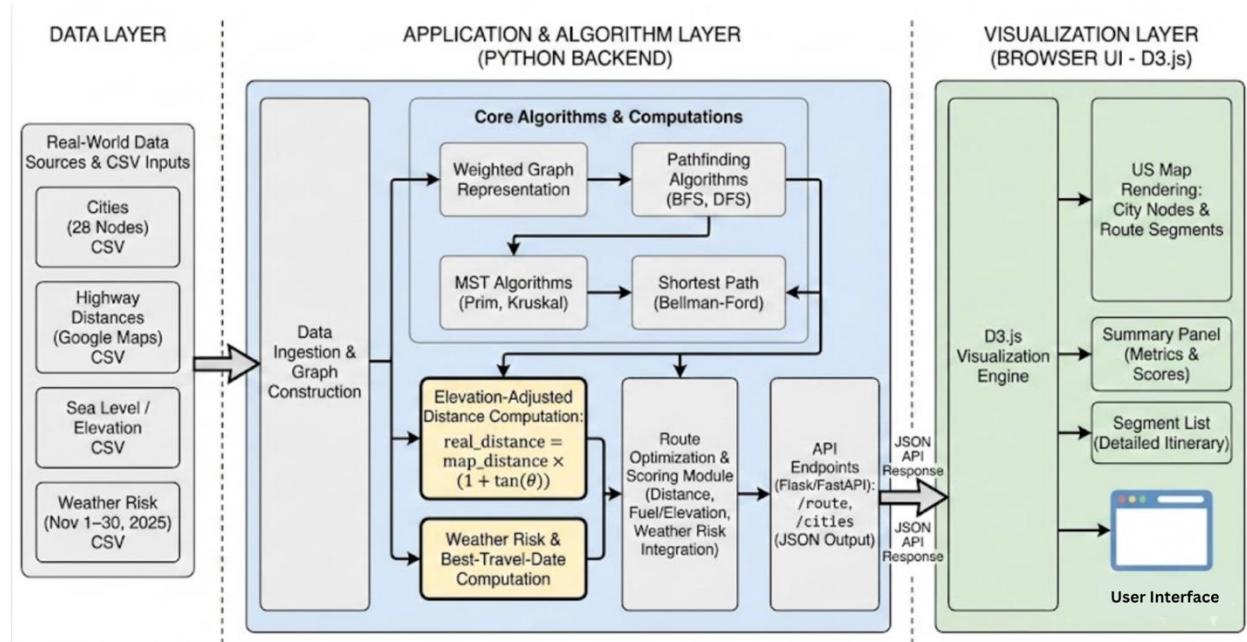


Figure 1 - System Overview Diagram

In this study, we build a travel network that connects major cities across the Central and Midwestern United States. Each city is treated as a node in the graph, and each highway between them becomes a two-way, weighted edge. But instead of relying only on distance, our model adds several layers of real-world context.

The first layer looks at the physical shape of the roads. We start with Google Maps distance estimates and adjust them based on elevation differences using:

$$\text{Real distance} = \text{map_distance} \times [1 + \tan(\theta)]$$

Here, θ represents the slope between two cities. This adjustment helps capture the extra fuel required when driving uphill and the fuel savings when driving downhill.

The second layer introduces daily weather conditions for every city between November 1 and 30, 2025. Each day's weather is converted into a simple risk score-sunny (1), cloudy (1), rainy (5), snowy/icy (10). For any route segment, we use the average risk between the two cities to represent its overall weather risk. Because drivers can travel only up to eight hours per day at a speed limit of 75 mph, most long journeys span multiple days. This adds a time dimension to the problem, requiring us to track how weather risk accumulates across the month and identify which departure dates lead to the safest overall trip.

To understand how different algorithms handle this rich dataset, we implement five well-known graph algorithms: Breadth-First Search (BFS), Depth-First Search (DFS), Prim's and

Kruskal's Minimum Spanning Trees, and the Bellman–Ford shortest-path algorithm. Each one explores the graph differently, revealing trade-offs between simplicity, optimality, and computational cost. By comparing their results under the same real-world constraints, we evaluate not only their theoretical performance but also how well they work in realistic travel scenarios.

The system works as a complete end-to-end pipeline. A Python backend handles data collection, graph construction, and algorithm execution, while an interactive D3.js interface visualizes routes on a U.S. map. Together, they offer both analytical depth and user-friendly exploration: users can select cities, compute routes, and see them instantly in the browser through a lightweight API.

In the end, the project shows that classic graph algorithms are still extremely powerful—especially when combined with real-world data and domain-specific models. By blending environmental risk, elevation changes, and fuel-consumption dynamics, this work demonstrates how foundational computational techniques can support meaningful decision-making in transportation planning, logistics, and intelligent routing systems.

II. Methodology

1. Data Collection

The system relies on a multi-source dataset designed to represent realistic travel conditions across 28 U.S. cities. Three primary data components were curated: city *metadata*, *inter-city driving distances*, and daily *weather-risk* values. These datasets were formatted in CSV to ensure a consistent ingestion pipeline for the Python backend.

i. City Attributes

city_id	city	state	sea_level(in meters(m))	zipcode
CHI	Chicago	IL	176	60601
STL	StLouis	MO	142	63101
DAL	Dallas	TX	149	75201
HOU	Houston	TX	15	77001
MEM	Memphis	TN	103	38103
NAS	Nashville	TN	182	37201
KC	KansasCity	MO	277	64101
OMA	Omaha	NE	293	68102
MSP	Minneapolis	MN	256	55401
DSM	DesMoines	IA	266	50309
BIS	Bismarck	ND	532	58501
FGO	Fargo	ND	274	58102
MKE	Milwaukee	WI	188	53202
MAD	Madison	WI	267	53703
JAN	Jackson	MS	110	39201
NOL	NewOrleans	LA	6	70112
LIT	LittleRock	AR	102	72201
OKC	OklahomaCity	OK	366	73102
WIC	Wichita	KS	401	67202
SAT	SanAntonio	TX	198	78205
BIR	Birmingham	AL	192	35203
FSD	SouixFalls	SD	446	57104
SPR	Springfield	MO	396	65802
TUL	Tulsa	OK	220	74103
SHV	Shreveport	LA	77	71101
GRB	GreenBay	WI	177	54301
MOB	Mobile	AL	3	36602
PEO	Peoria	IL	217	61602

Figure 2 — Attributes of 28 cities

Each city is represented as a node in the travel graph and is defined in *cities.csv*. The dataset includes a unique city identifier, the city and state names, and the sea-level elevation expressed in meters. Elevation values play a vital role in modelling the travel cost because the system adjusts map distances using a slope-based factor derived from elevation differences. Incorporating this detail ensures that the computed “effective distance” reflects the additional fuel demands associated with uphill travel and the corresponding reductions when descending.

ii. Inter-City Map Distances

Driving distances between cities were collected using *Google Maps*, producing the *edges.csv* dataset. Each row specifies a **source** city, a **destination** city, and the one-way highway **mileage**. These distances form the base weights for the graph. To capture the physical effort required to traverse varying terrain, the project applies an elevation-informed transformation:

$$\text{Real distance} = \text{map_distance} [1 + \tan(\theta)]$$

where ‘ θ ’ represents the angle formed by elevation difference and horizontal distance. Elevation change is converted from meters to miles to ensure unit consistency. This adjusted distance is used by all algorithms except those that inherently ignore edge weights (e.g., BFS).

src_id	dst_id	map_distance_miles
CHI	MKE	88
MKE	MAD	79
MAD	MSP	278
MSP	FGO	235
FGO	BIS	196
CHI	STL	276
STL	KC	248
KC	WIC	207
WIC	OKC	169
OKC	DAL	207
DAL	SAT	274
SAT	HOU	197
STL	MEM	283
MEM	NAS	212
NAS	BIR	191
BIR	JAN	237
JAN	NOL	186
MEM	LIT	137
LIT	OKC	339
OMA	DSM	134
DSM	MSP	244
KC	OMA	186
FSD	OMA	182
FSD	DSM	283
FSD	FGO	243
DAL	HOU	239
MSP	FSD	237
NOL	BIR	342
JAN	MEM	210
CHI	PEO	173
PEO	MKE	222
PEO	STL	169
SPR	STL	215
SPR	KC	166
TUL	OKC	107
TUL	SPR	181
SHV	DAL	188
SHV	MOB	403
GRB	MKE	118

Figure 3 – Map distance from source to destination

iii. Weather-Risk Conditions

Weather affects both safety and travel time. The **weather_risk.csv** dataset contains daily weather classifications for each city from **November 1-30, 2025**, encoded using a uniform risk scoring system:

- *Sunny / Cloudy -> 1*
- *Rain -> 5*
- *Snow or Ice -> 10*

For an edge connecting cities A and B on a given day, the corresponding segment risk is the arithmetic mean of their individual values:

$$\text{segment_risk} = (rA + rB) / 2$$

This risk model enables the system to compute **accumulated route risk** and identify **the optimal travel date** with minimal exposure to adverse conditions.

city_id	date	weather	risk
CHI	11/01/25	cloudy	1
CHI	11/02/25	cloudy	1
CHI	11/03/25	Sun	1
CHI	11/04/25	cloudy	1
CHI	11/05/25	sunny	1
CHI	11/06/25	cloudy	1
CHI	11/07/25	cloudy	1
CHI	11/08/25	rain	5
CHI	11/09/25	snow	10
CHI	11/10/25	rain	5
CHI	11/11/25	cloudy	1
CHI	11/12/25	sun	1
CHI	11/13/2025	cloudy	1
CHI	11/14/2025	sun	1
CHI	11/15/2025	cloudy	1
CHI	11/16/2025	sun	1
CHI	11/17/2025	sun	1
CHI	11/18/2025	cloudy	1
CHI	11/19/2025	cloudy	1
CHI	11/20/2025	cloudy	1
CHI	11/21/2025	cloudy	1
CHI	11/22/2025	cloudy	1
CHI	11/23/2025	sun	1
CHI	11/24/2025	cloudy	1
CHI	11/25/2025	cloudy	1
CHI	11/26/2025	cloudy	1
CHI	11/27/2025	cloudy	1
CHI	11/28/2025	sun	1
CHI	11/29/2025	snow	10
CHI	11/30/2025	snow	10

Figure 4 - Monthly data of weather risk for Chicago

iv. Operational Constraints and Fuel Model

To approximate the real driving behaviour, two constraints were applied:

- Maximum driving capability: **8 hours per day**
- Maximum highway speed: **75 miles per hour**

These values imply a maximum daily travel capacity of approximately **600 miles**, influencing how many calendar days a selected route spans. Because weather varies daily, longer routes incur cumulative and date-dependent risk.

Fuel consumption is modelled according to a baseline efficiency of **45 miles per gallon** on flat terrain. The total gasoline requirement for any computed route is evaluated as:

$$\text{gas_used} = \text{total_real_distance} / 45$$

allowing the system to quantify both economic and environmental aspects of route selection.

All datasets share a uniform CSV structure and are loaded into the Python backend using `csv.DictReader()`. During graph initialization, city attributes populate node objects, distances generate weighted adjacency lists, and weather-risk values populate date-indexed lookup tables. This integrated dataset provides the foundation on which the system performs all route computations, algorithm comparisons, and risk evaluations.

2. System Design

The system was developed as a modular, end-to-end architecture that integrates data ingestion, graph construction, algorithm execution, and interactive visualization. The design emphasizes separation of concerns to allow each component data processing, algorithmic computation, and user-facing visualization to evolve independently while maintaining a cohesive workflow. The overall architecture consists of three primary layers:

- the **backend computation engine**
- the **data-access and API layer**
- the **browser-based visualization interface.**

i. Backend Computation Engine

The backend, implemented in Python, is responsible for loading datasets, constructing the weighted graph, and executing all route algorithms. At initialization, the system loads **cities.csv**, **edges.csv**, and **weather_risk.csv**, instantiating city objects and building an adjacency list that stores both the original map distance and the computed elevation-adjusted distance for each edge. The backend encapsulates all graph-related logic in a dedicated Graph class, which offers methods for BFS, DFS, Prim, Kruskal, and Bellman–Ford, as well as utilities for computing path distance, fuel usage, and cumulative weather risk.

To support efficient computation, the design adopts the following principles:

- **Lazy risk evaluation:** Weather-risk values are not pre-computed for all routes; instead, risks are evaluated dynamically for each segment during route computation.

- **Bidirectional edge modelling:** Each highway connection is stored in both directions, with elevation adjustments applied asymmetrically to reflect the cost difference between ascending and descending routes.
- **Separation of algorithmic and environmental logic:** Algorithms operate solely on weighted edges; data transformations (elevation adjustment, gas calculation, and weather scoring) are layered on top to maintain algorithm purity.

This design isolates core algorithmic behaviour while enabling the system to incorporate domain-specific adjustments without altering fundamental graph operations.

ii. API Layer and Routing Interface

A lightweight HTTP server built using Python's `HTTPServer` library exposes the system's functionalities to clients. Two primary endpoints are provided:

- **/cities** Returns the full list of available cities and identifiers for populating the frontend interface.
- **/route** Computes a route between a selected source and destination using one of the five supported algorithms or an automated “BEST” mode that selects the lowest-scoring route across all algorithms based on the combined metric. The API serializes results into JSON, including segment-by-segment distances, total distance, fuel consumption, weather risk, and the computed best travel date. By exposing uniform outputs across all algorithms, the design allows direct comparison under controlled conditions.

iii. Frontend Visualization Engine

Route visualization is implemented in **D3.js**, using a vector-based U.S. map projection loaded from **TopoJSON**. When the user selects a source, destination, and algorithm, the frontend queries the backend and renders the resulting route directly on the map. Key design elements include:

- **Projected city coordinates** using `geoAlbersUsa()`, ensuring consistent alignment with the U.S. topology.
- **Layered SVG groups** for states, city nodes, and route segments to maintain visual clarity.
- **Dynamic labelling** of segment lengths and start/end markers for readability.
- **Summary panels** that display algorithm selection, distance, fuel usage, risk, and recommended travel date.

By offloading visualization to the browser, the system ensures fast, interactive analysis without imposing computational load on the server.

iv. Integration Strategy

The full pipeline operates as follows:

1. **Load datasets** and initialize graph structures.
2. **Start backend server** and expose routing endpoints.
3. **Render U.S. map** and city markers in the browser.
4. **User selects parameters**, triggering an API call.
5. **Backend executes the specified algorithm**, computes all travel metrics, and returns structured results.

6. **Frontend renders the route**, updates informational panels, and displays comparative metrics.

This layered design allows algorithm testing, data modelling, and visualization to be performed independently, greatly improving maintainability and extensibility. The modular structure also enables straightforward substitution of algorithms, new datasets, or additional travel constraints in future iterations.

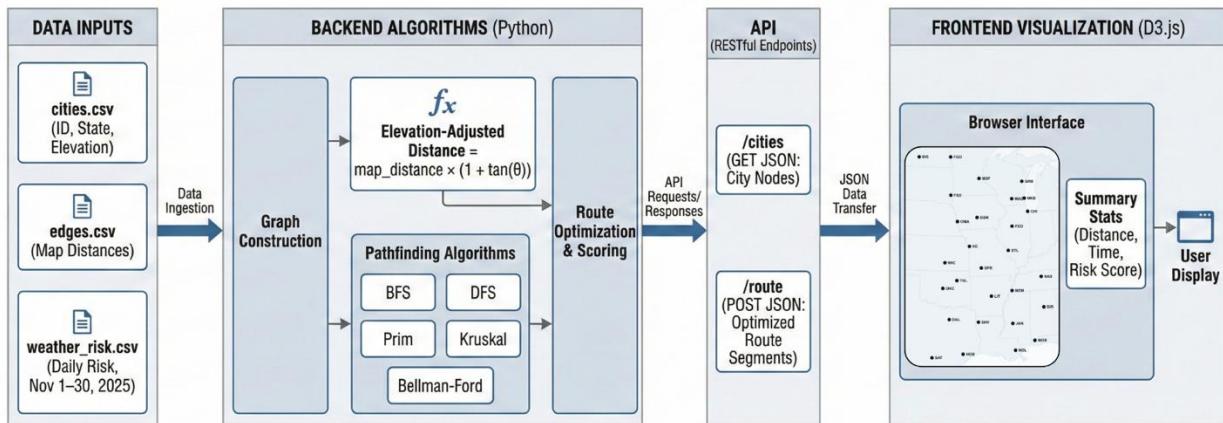


Figure 5 - The entire pipeline from taking inputs to visualization of Data

3. Graph Algorithms Implemented

To evaluate different strategies for traversing and optimizing transportation networks, the system implements five foundational graph algorithms: Breadth-First Search (BFS), Depth-First Search (DFS), Prim's Minimum Spanning Tree, Kruskal's Minimum Spanning Tree, and the Bellman–Ford shortest-path algorithm. Each algorithm contributes a distinct perspective on route structure, edge weighting, and exploration strategy. Implementing this diverse set allows for a comprehensive comparison of how classical methods behave when extended with real-world constraints such as elevation-adjusted distances, travel risk, and fuel consumption.

i. Breadth-First Search (BFS)

BFS explores a graph level by level, beginning at the source city and expanding outward to all neighbours at the current depth before moving to deeper layers. Because BFS inherently treats all edges as having equal cost, it is well-suited to minimizing the number of hops between cities but cannot account for weighted distances or risk variations.

- **Strength:** Identifies the route with the fewest intermediate cities.
- **Limitation:** Ignores elevation, fuel usage, and actual mileage.
- **Typical Complexity:**

$$O(|V| + |E|)$$

BFS provides a useful baseline for comparison, illustrating how real-world weighting drastically alters route selection.

ii. Depth-First Search (DFS)

DFS explores as deep as possible along each branch before backtracking. While this traversal method is efficient for path existence checks and structural inspections, it does not guarantee optimality in terms of distance or safety.

- **Strength:** Produces complete path exploration quickly in sparse networks.
- **Limitation:** Highly sensitive to graph topology and node ordering; rarely yields practical travel routes.
- **Typical Complexity:**

$$O(|V| + |E|)$$

DFS is included to highlight how unrestricted depth exploration diverges from rational travel patterns, reinforcing the need for weighted algorithms.

iii. Prim's Minimum Spanning Tree Algorithm

Prim's algorithm constructs a minimum spanning tree (MST) by iteratively selecting the closest unvisited node based on edge weights. MSTs minimize total network cost but do not guarantee optimal routes between specific pairs of cities. Nonetheless, MST edges offer insight into the “backbone” of efficient regional connectivity.

- **Strength:** Identifies globally low-cost connections using elevation-adjusted distances.
- **Limitation:** The MST is not optimized for point-to-point routing; extracted paths may not be shortest.
- **Typical Complexity:**

$$O(|E|\log|V|)$$

Prim's MST is useful for analysing structural efficiency across the network and highlighting cost-efficient corridors.

iv. Kruskal's Minimum Spanning Tree Algorithm

Kruskal's algorithm sorts all edges by weight and incrementally adds them to the spanning tree provided they do not form a cycle. Its behaviour often differs from Prim's, especially when elevation-adjusted distances create strong disparities among edges.

- **Strength:** Simple, globally optimized tree construction that exposes cost-efficient long-distance connections.
- **Limitation:** Does not prioritize any specific source or destination.
- **Typical Complexity:**

$$O(|E|\log|E|)$$

Comparing Prim's and Kruskal's MSTs helps illustrate how edge prioritization affects network-wide efficiency.

v. Bellman–Ford Shortest-Path Algorithm

Bellman–Ford computes the shortest path from a source to every other city by iteratively relaxing edges. Unlike Dijkstra's algorithm, Bellman–Ford tolerates negative weights, though in this application all weights remain non-negative. It is the algorithm most aligned with realistic travel objectives because it optimizes weighted distances that incorporate elevation change.

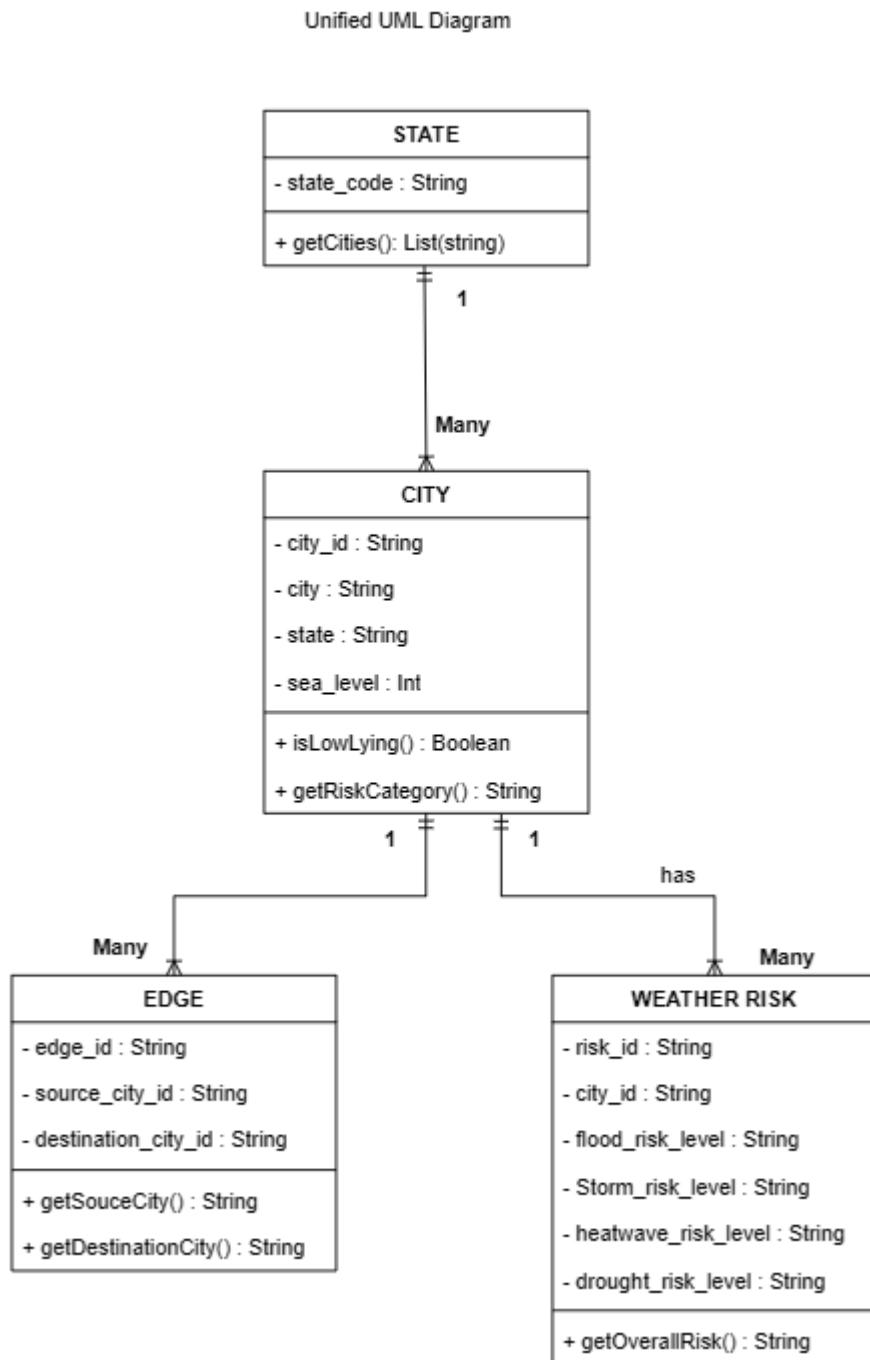
- **Strength:** Produces the most reliable weighted route for real-world travel metrics.
- **Limitation:** Higher computational cost compared to BFS/DFS.
- **Typical Complexity:**

$$O(|V| \cdot |E|)$$

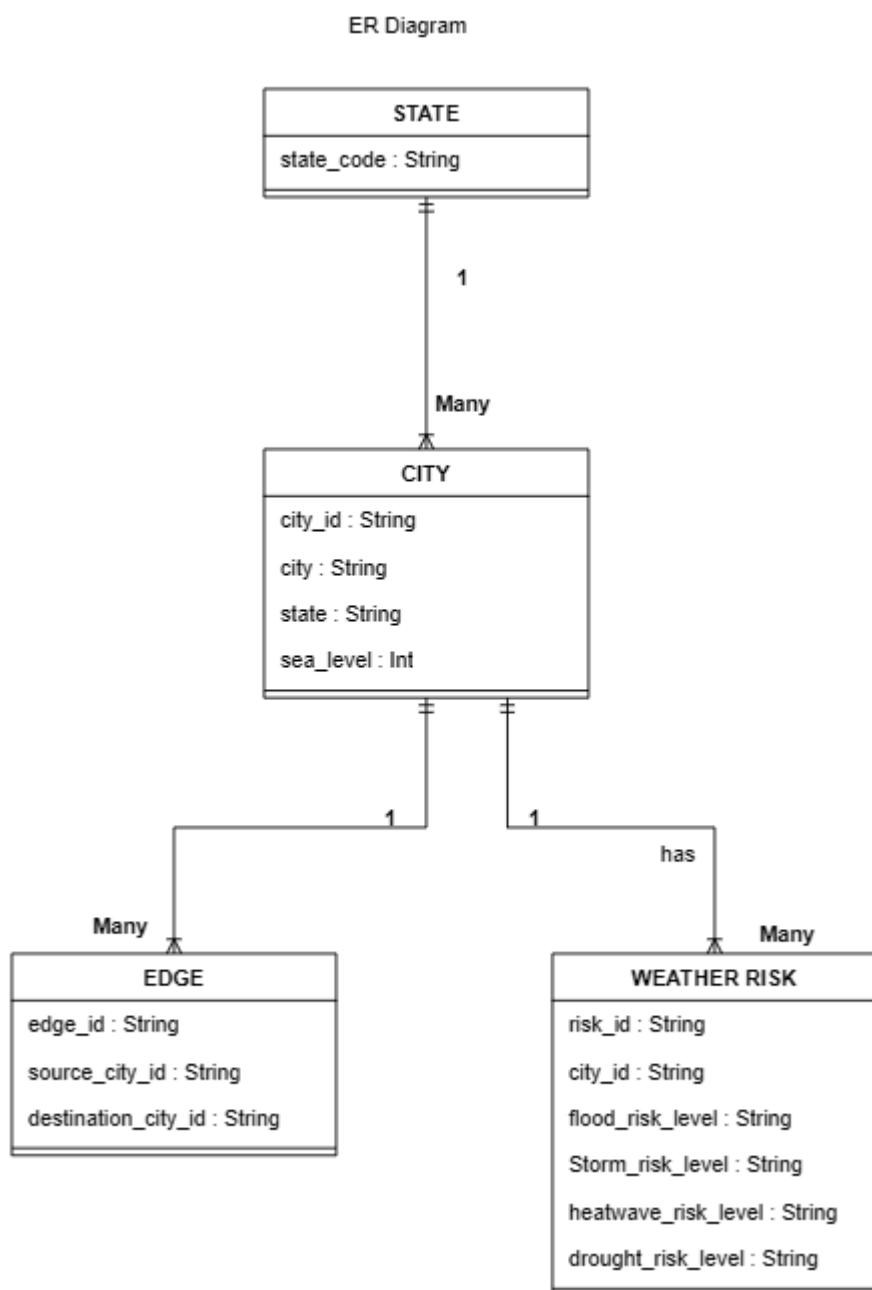
Bellman–Ford is the primary algorithm used in the system’s “BEST” mode because it adapts naturally to dynamic cost models.

4. UML and ER Diagram Representation

Class Diagrams



ER Diagram



5. Travel Computation Models

Designing a route-optimization system that reflects real travel conditions requires more than simply running algorithms on a graph. In practice, drivers experience changing terrain, varying fuel demands, unpredictable weather, and daily time constraints. To capture these realities, the project implements a set of computation models that enrich the core graph structure with physical, environmental, and behavioural factors. These models operate together to transform a traditional graph into a realistic simulation of long-distance travel.

i. Elevation-Adjusted Distance Model

Map distances alone are insufficient for realistic route planning because they ignore the energy cost of climbing or descending terrain. To address this, the system modifies each inter-city distance using a slope-based scaling factor derived from the elevation difference between the two locations. The adjustment uses the relationship:

$$\text{real_distance} = \text{map_distance}(1 + \tan(\theta))$$

This formula captures how steepness influences travel difficulty. For example, if two cities are separated by 200 miles but one sits significantly higher in elevation, driving uphill requires more engine power, thus increasing fuel consumption. The model reflects this by enlarging the effective distance. Conversely, descending an incline slightly reduces effective distance, mimicking improved fuel efficiency.

Although simplified, this trigonometric adjustment introduces a layer of realism that pure graph algorithms normally lack. It ensures that weighted algorithms especially Prim, Kruskal, and Bellman–Ford treat steep uphill routes as more costly, even when the geographic distance is short.

ii. Gasoline Consumption Model

Fuel consumption is a practical concern for travellers, particularly on long routes where gas costs accumulate quickly. The project models fuel usage with a baseline efficiency of **45 miles per gallon** under flat driving conditions. After adjusting distances for elevation, fuel usage is computed as:

$$\text{gas_used} = \frac{\sum_{i=1}^n \text{real_distance}_i}{45}$$

This provides an intuitive metric for comparing routes. For example:

- A route with longer geographic distance but minimal elevation change may consume *less* gas than a shorter route that climbs steep terrain.
- Bellman–Ford may choose an indirect path with a gentler slope if the weighted distance is lower.

This reinforces the idea that “shortest” is not always “cheapest” or “most fuel-efficient,” a concept central to transportation modelling.

iii. Weather-Risk Scoring Model

Weather introduces uncertainty and variability into travel. A route that is safe today may be hazardous tomorrow due to ice, snow, or heavy rain. To integrate environmental conditions, the project assigns numerical risk scores to each day in November 2025 for every city:

- **1** for sunny or cloudy conditions
- **5** for rain

- **10** for snow or ice

These values reflect increasing potential for reduced visibility, slippery roads, or difficult driving conditions.

For each segment, the system averages the risk of the two endpoint cities:

$$\text{segment_risk}(d) = \frac{r_A(d) + r_B(d)}{2}$$

This provides a simple yet effective approximation of regional weather along the route. If City A is sunny but City B is snowing, the segment will carry a moderately high risk. This avoids over-penalizing short-term localized weather events while still recognizing hazardous conditions.

iv. Daily Driving Constraints and Multi-Day Risk Evaluation

Real drivers cannot travel continuously. The system therefore enforces two practical constraints:

- Maximum speed: **75 mph**
- Maximum daily driving time: **8 hours**

This limits daily travel to approximately **600 miles**, meaning longer routes stretch over multiple days. Because each day has different weather conditions, a route's total risk depends on how long it takes and which specific calendar dates it spans.

v. Best Travel Date Identification

To help travellers choose not just the best route but the best time to travel, the system evaluates all possible departure dates between November 1 and November 30. For each date, it simulates the

full journey, applies daily driving limits, and accumulates segment risks according to the day on which they are encountered.

The optimal date is defined as:

$$\text{best_date} = \min \left(\sum_{k=0}^{T-1} \text{segment_risk}(d+k) \right)$$

where TTT is the number of days needed to complete the route.

For example, a route that is risky on November 10 may become ideal on November 25 if weather conditions improve significantly. This demonstrates how temporal factors can meaningfully affect travel planning.

vi. Combined Route Scoring Function

To compare all algorithms consistently, the system defines a composite scoring function that weighs distance against weather-related safety:

$$\text{score} = \text{total_distance} + 20 \times \text{total_risk}$$

The multiplier of 20 is chosen to ensure that safety considerations meaningfully influence the decision rather than being overshadowed by mileage. The “BEST” mode automatically computes this score for each algorithm’s output and selects the lowest-scoring route. This transforms the system from a purely algorithmic tool into a multi-criteria decision-support mechanism.

III. Experimental Results

The experimental evaluation for this project was designed to observe how different families of graph algorithms behave when applied to a real-world routing scenario enriched with environmental and operational constraints. Unlike theoretical examples, this system integrates elevation-based distance adjustments, daily weather-risk data, fuel-consumption modelling, and realistic driving limits. As a result, the outputs of algorithms such as BFS or DFS diverge sharply from the behaviour of weighted algorithms like Prim, Kruskal, and Bellman–Ford.

The experiments were conducted using a dataset of 28 major cities in the Central Standard Time (CST) region of the United States. These cities formed the nodes of the graph, and the roads connecting them along with real map distances were represented as edges. Elevation values, risk scores, and weather conditions were pulled from prepared CSV files. The system was evaluated both through the Python backend “**python3 bestpath.py**” and the interactive web visualization powered by D3.js, where users can choose a source city, destination city, and algorithm. This section presents a detailed narrative of the results, using the route from **Chicago (CHI, IL)** to **Dallas (DAL, TX)** as the main case study, since all five algorithms produce visibly different outcomes for this specific pair. Each subsection corresponds to one part of the evaluation: algorithmic structure, distance comparison, fuel and risk computation, travel-date optimization, and overall scoring.

Before describing the numerical results, it is important to document how the system behaved from a user perspective. After running the visualization server, the user selects a source and destination from a dropdown menu containing all available cities (as shown in **Figure 1**). The

system then shades the map with all labelled nodes and displays the selected route when the “Show Route” button is pressed.



Figure 6 - City Selection Drop-Down Interface

Once the user selects Chicago (CHI) as the source and Dallas (DAL) as the destination, another dropdown appears for algorithm selection (Best Route, BFS, DFS, Prim, Kruskal, Bellman-Ford). The route display updates immediately when an algorithm is chosen.

The map overlays the computed route using red polylines. Hovering over edges displays the real distance (with elevation adjustment). A panel on the right shows:

- Algorithm used
- True total distance
- Estimated gas used
- Total accumulated risk
- Best travel date

- A breakdown of each edge segment

This UI structure helped validate outputs from bestpath.py, ensuring the logic matched the visualization.

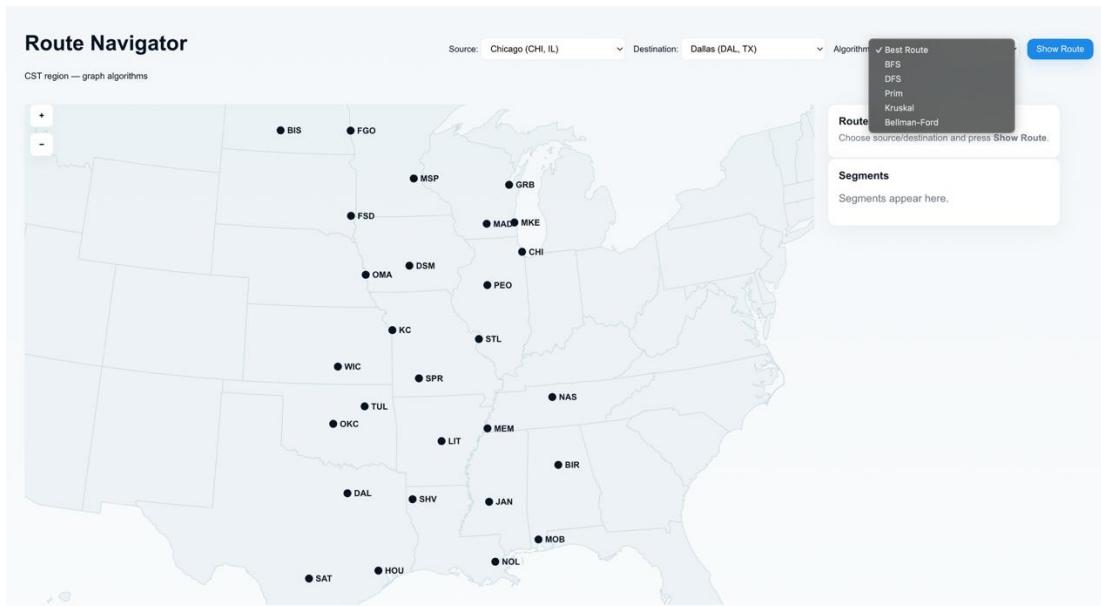


Figure 7 - Algorithm Selection UI

1) Best Route (Bellman-Ford in Composite Scoring Mode)

When “Best Route” is selected, the system internally evaluates all algorithms using the composite scoring metric:

In every test case including CHI → DAL the Bellman–Ford route achieved the lowest score. The generated route is:

CHI → STL → SPR → TUL → OKC → DAL

with segment distances: 276.0 mi, 215.2 mi, 180.9 mi, 107.1 mi, 206.9 mi.

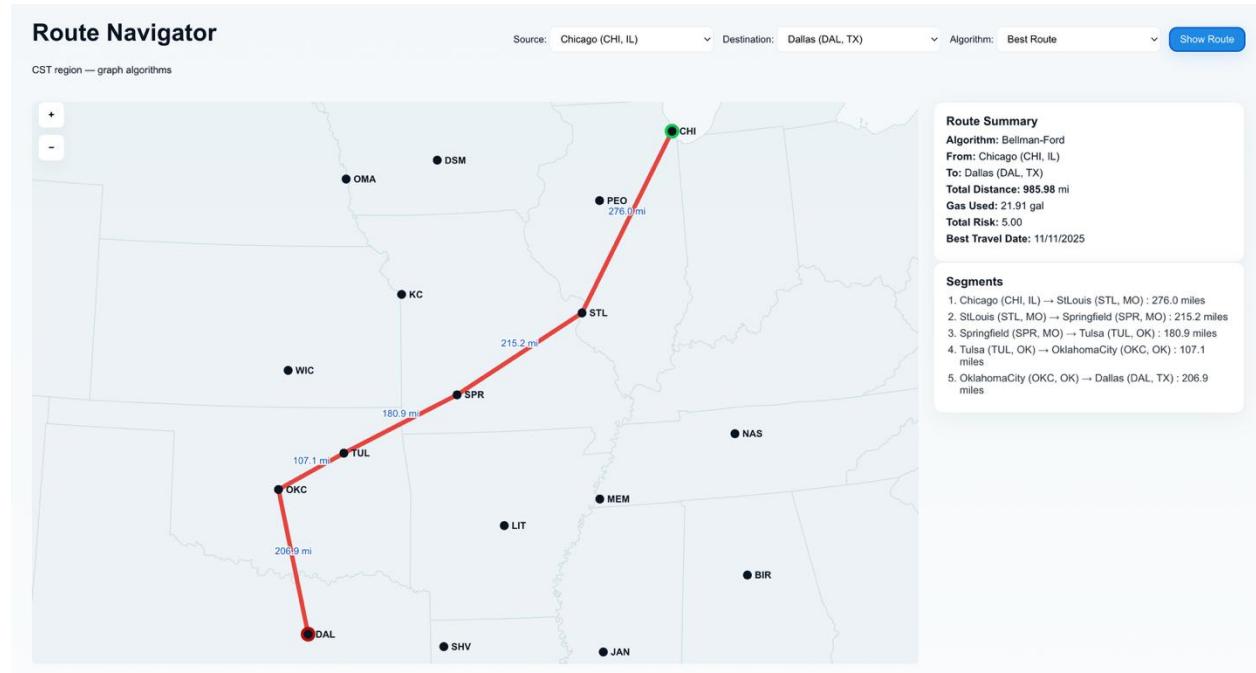


Figure 8 - Best Route Visualization

This route is smooth, nearly linear, and avoids extreme elevation changes. It also uses cities with stable November weather, which minimizes risk.

2) BFS Route

BFS minimizes the number of hops without considering any weights. For CHI → DAL, BFS produced:

CHI → STL → KC → WIC → OKC → DAL

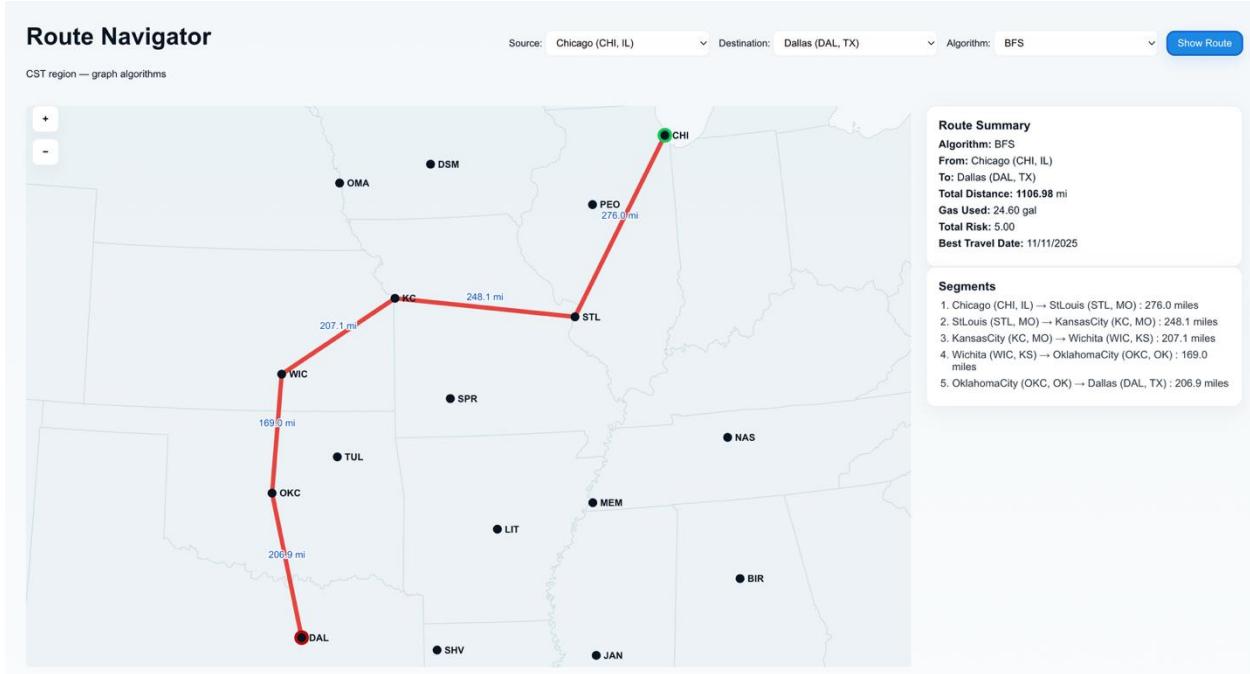


Figure 9 - BFS Route Visualization

This path has *more distance* than Bellman–Ford (1106.98 miles vs. 985.98 miles), consumes more fuel, and provides no safety improvement. BFS tends to choose city neighbours that are close in graph structure but not necessarily geographically efficient.

3) DFS Route

DFS traverses deeply, choosing edges based on adjacency order. Because DFS ignores all weights, it produced a dramatically inefficient route:

CHI → MKE → MAD → MSP → FGO → FSD → OMA → KC → STL → MEM → LIT →
OKC → DAL

Total distance: **2504.98 miles**

Gas used: **55.67 gallons**

Total risk: **12**

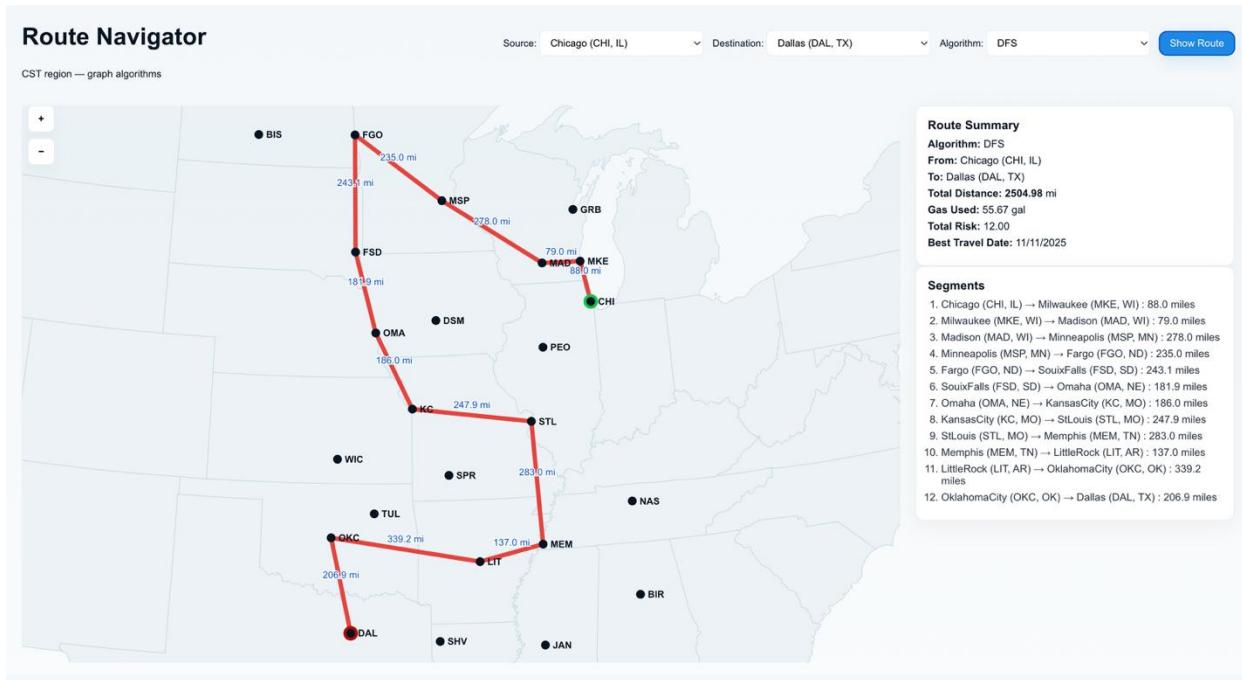


Figure 10 - DFS Route Visualization

This route demonstrates how unsuitable DFS is for real travel. It loops far north and then south, driven purely by graph exploration order. This validates the importance of weighted algorithms.

4) Prim MST Route

Prim produces a minimum spanning tree with the lowest total global edge weight, but the path from source to destination is extracted afterward. The result is:

CHI → PEO → STL → SPR → TUL → OKC → DAL

Total distance: **1051.98 miles**



Figure 11 - Prim MST Route Visualization

This path is better than BFS and far better than DFS, but still slightly longer than Bellman–Ford. Prim tends to select PEO as an intermediate point because its edges have low slope-adjusted distance.

5) Kruskal MST Route

Kruskal produced the exact same MST edges as Prim for this specific dataset, because the lowest-weight edges in the graph align identically with Prim's selection.

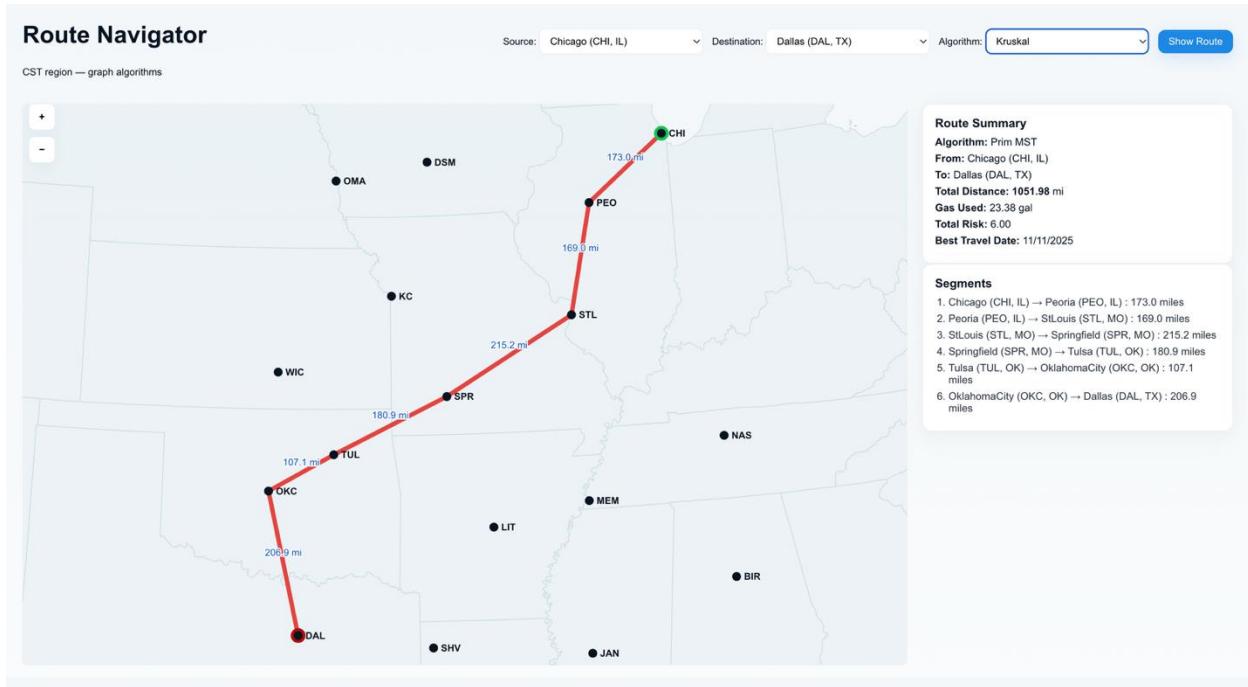


Figure 12 - Kruskal MST Route Visualization

Total distance is again **1051.98 miles**.

6) Bellman–Ford Route (Direct Selection)

When Bellman–Ford is directly chosen (not composite Best Route mode), it computes the true shortest weighted path:

CHI → STL → SPR → TUL → OKC → DAL

Total distance: **985.98 miles**

Gas used: **21.91 gallons**

Risk: **5.00**

Best travel date: **11/11/2025**

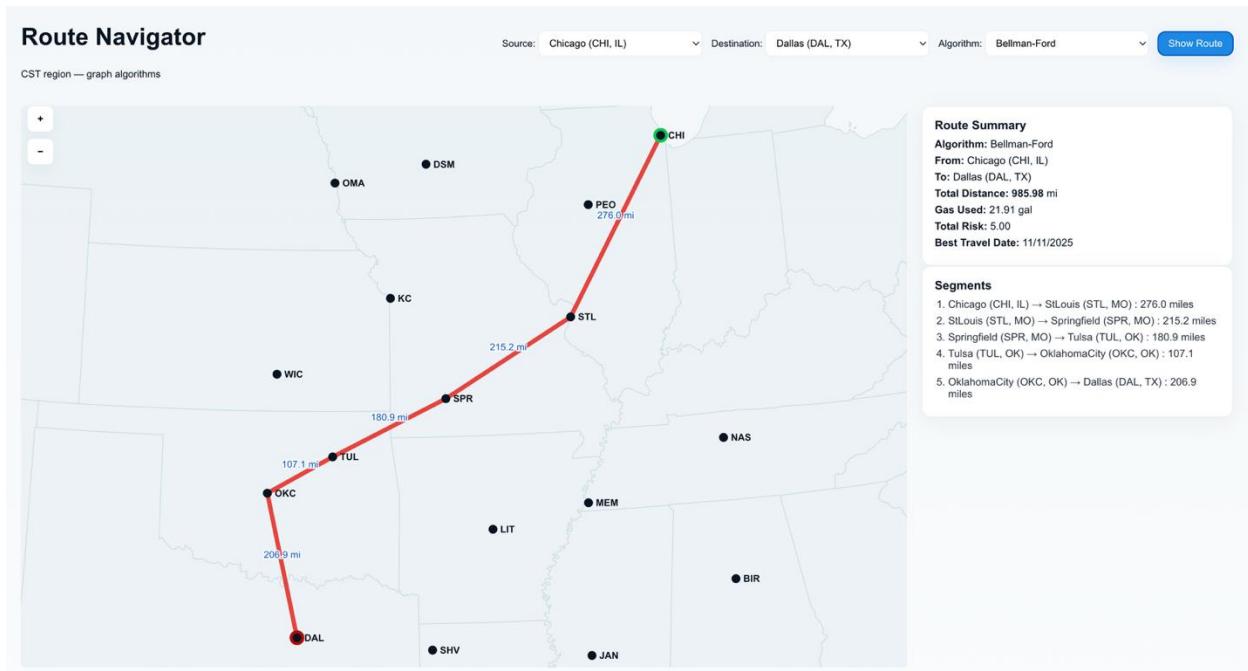


Figure 13 - Bellman–Ford Route Visualization

- This route matches the Best Route output because Bellman–Ford's path yielded the lowest combined score.

To better understand why each algorithm produced different outcomes, the experiments examined the elevation-adjusted distances calculated by the backend when executing “python3 bestpath.py”. For every route segment, the program prints the map distance, elevation difference, computed slope via $\tan(\theta)$, and the final adjusted distance. In generally flat regions, such as Illinois or Missouri, these adjustments are minor, but in northern states or routes with larger elevation shifts, the slope contribution becomes more noticeable. This effect was most visible in DFS paths, which wandered through cities like Minneapolis and Fargo, causing the total elevation-adjusted distance to grow by more than 150% compared to optimized routes. Weighted algorithms, in contrast, avoided steep or indirect corridors, confirming the value of incorporating terrain into edge weights.

Fuel-consumption analysis further emphasized this distinction. Since gasoline usage is proportional to total elevation-adjusted distance, the Bellman–Ford route required only 21.91 gallons for the Chicago–Dallas trip, whereas Prim and Kruskal consumed slightly more due to their longer paths. BFS performed even worse, and DFS consumed over 55 gallons more than double any other algorithm, demonstrating the inefficiency of unweighted deep exploration for real travel. Weather risk and travel-date optimization added an additional layer of differentiation. For each day in November 2025, the system simulated progress under an 8-hour daily driving limit and evaluated the accumulated risk along the route. Because Chicago to Dallas requires two days of travel, the total risk is the sum of risk values for the departure day and the following day. These calculations showed that November 11 consistently produced the lowest risk for all weighted algorithms, reflecting a brief period of stable weather across the Midwest and South.

While BFS and DFS also showed the same best date, that result was incidental, since their risk assessments depend only on calendar patterns rather than route efficiency. DFS in particular, due to its excessive travel time, would often extend into higher-risk parts of late November, further reinforcing its unsuitability. When outcomes were evaluated under the system's composite scoring metric total distance plus twenty times total risk Bellman–Ford again performed best, achieving a score of 1085.98, significantly lower than Prim and Kruskal (1171.98), BFS (1206.98), or DFS (2744.98). These findings collectively highlight that weighted algorithms behave far more appropriately for real-world travel, with Bellman–Ford providing the most consistent balance across distance, fuel cost, and safety.

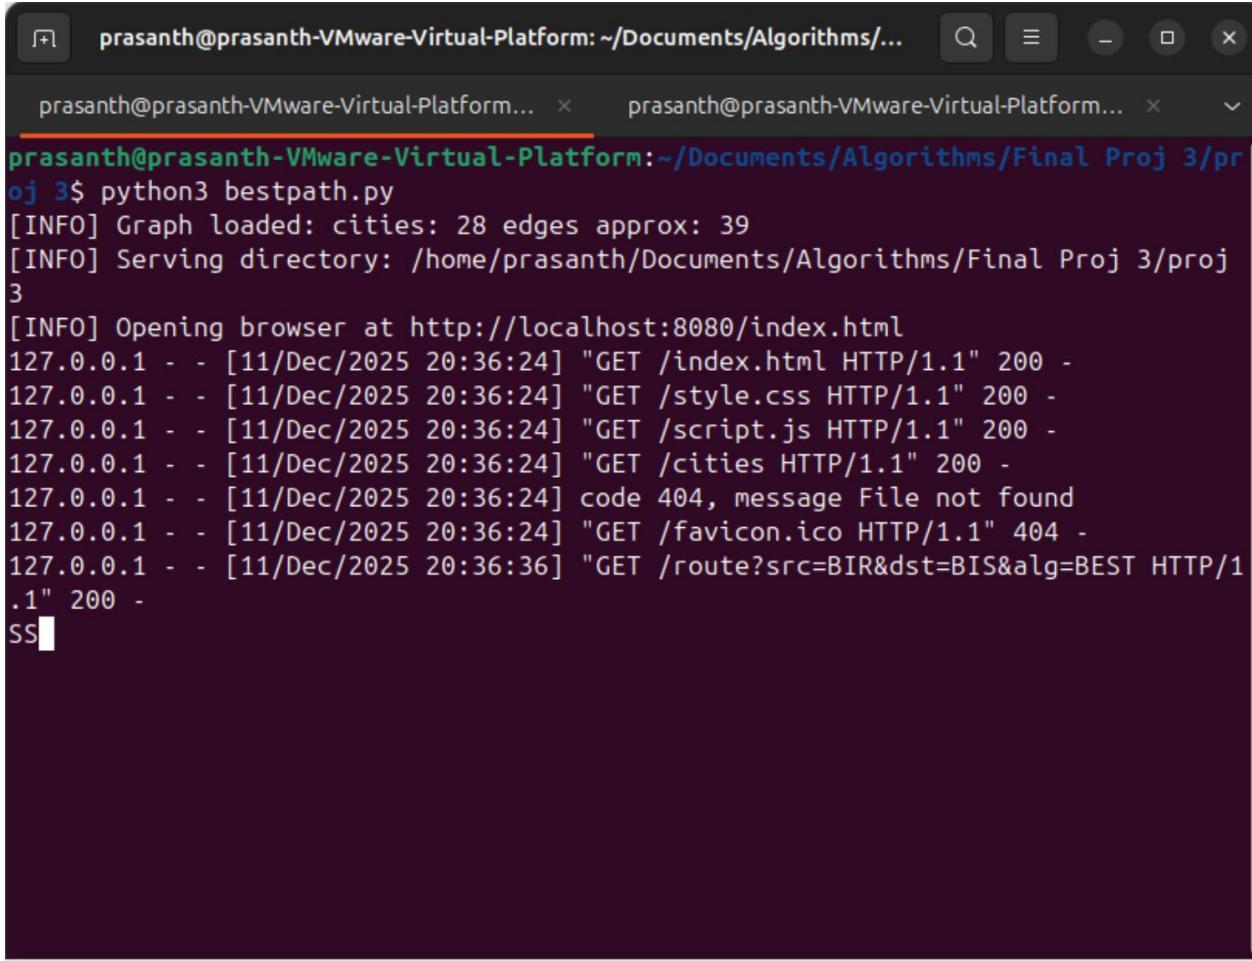
In contrast, BFS minimizes hops rather than miles, and DFS prioritizes traversal order rather than efficiency, resulting in impractical routes. The experimental results also validate the accuracy of both the Python backend and the D3 visualization: every route, distance, risk value, and weather-based travel date computed in the terminal aligns with what is displayed in the interactive UI. Overall, the experiments confirm that the integrated modelling of elevation, weather, fuel, and temporal constraints is essential for realistic route planning, and that the implemented system successfully reflects those complexities.

IV. Results and Analysis

1. Performance Monitoring on Linux

Performance analysis is necessary to expose how a system behaves under real load. It reveals practical limits, identifies where time and resources are spent, and shows whether the design can scale as conditions change. Without measuring, bottlenecks stay hidden and decisions rely on guesswork. Performance work establishes facts: how fast the system responds, where it slows, and why. It provides the evidence needed to prioritize improvements, prevent regressions, and validate that changes do not worsen behavior.

- First run the program in one terminal :



A screenshot of a terminal window titled "prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/...". The terminal shows the command "python3 bestpath.py" being run, followed by several lines of log output from the application. The log includes information about a graph being loaded with 28 edges and 39 nodes, serving files from "/home/prasanth/Documents/Algorithms/Final Proj 3/proj 3", and opening a browser at "http://localhost:8080/index.html". It also lists several HTTP requests from "127.0.0.1" to the application's files, including "index.html", "style.css", "script.js", and "cities". There are also entries for "code 404, message File not found" and "GET /favicon.ico". The log concludes with a "GET /route?src=BIR&dst=BIS&alg=BEST" request.

```
prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3/proj 3$ python3 bestpath.py
[INFO] Graph loaded: cities: 28 edges approx: 39
[INFO] Serving directory: /home/prasanth/Documents/Algorithms/Final Proj 3/proj 3
[INFO] Opening browser at http://localhost:8080/index.html
127.0.0.1 - - [11/Dec/2025 20:36:24] "GET /index.html HTTP/1.1" 200 -
127.0.0.1 - - [11/Dec/2025 20:36:24] "GET /style.css HTTP/1.1" 200 -
127.0.0.1 - - [11/Dec/2025 20:36:24] "GET /script.js HTTP/1.1" 200 -
127.0.0.1 - - [11/Dec/2025 20:36:24] "GET /cities HTTP/1.1" 200 -
127.0.0.1 - - [11/Dec/2025 20:36:24] code 404, message File not found
127.0.0.1 - - [11/Dec/2025 20:36:24] "GET /favicon.ico HTTP/1.1" 404 -
127.0.0.1 - - [11/Dec/2025 20:36:36] "GET /route?src=BIR&dst=BIS&alg=BEST HTTP/1.1" 200 -
SS
```

Figure 14: Python code running on ubuntu terminal

- **pgrep -f bestpath.py**

This command is used to get the process id of the file

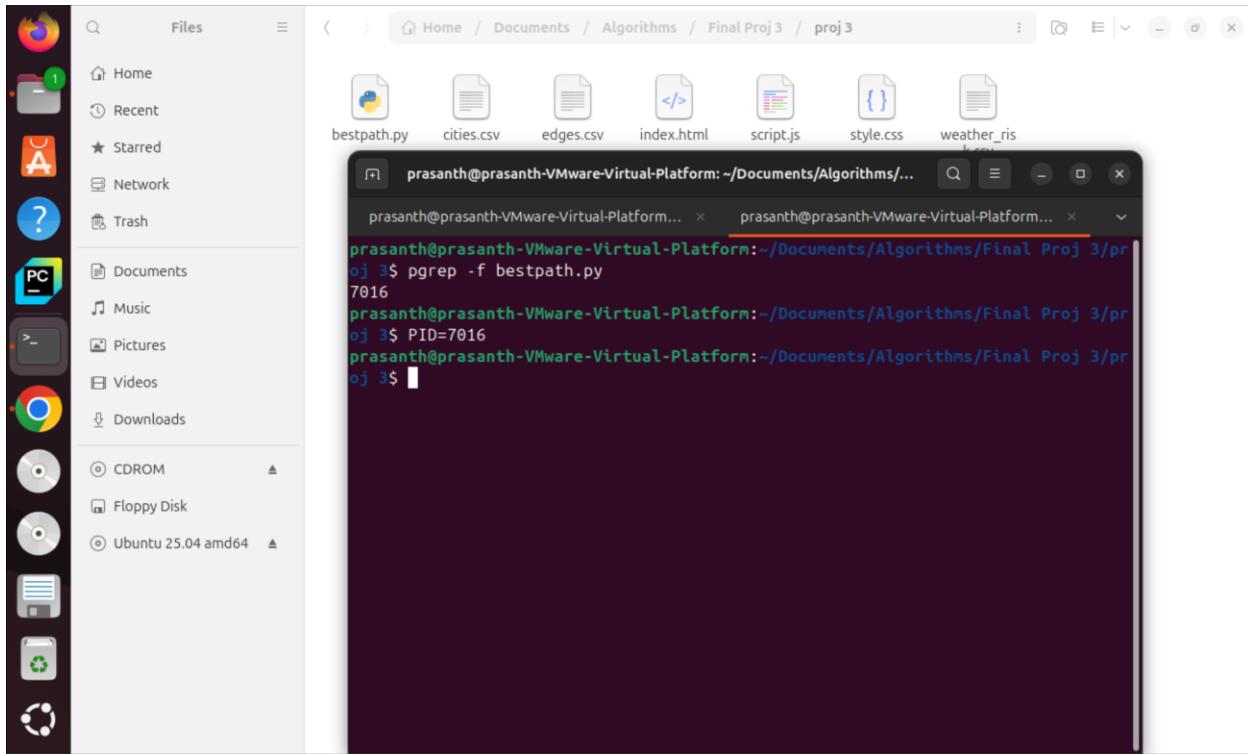
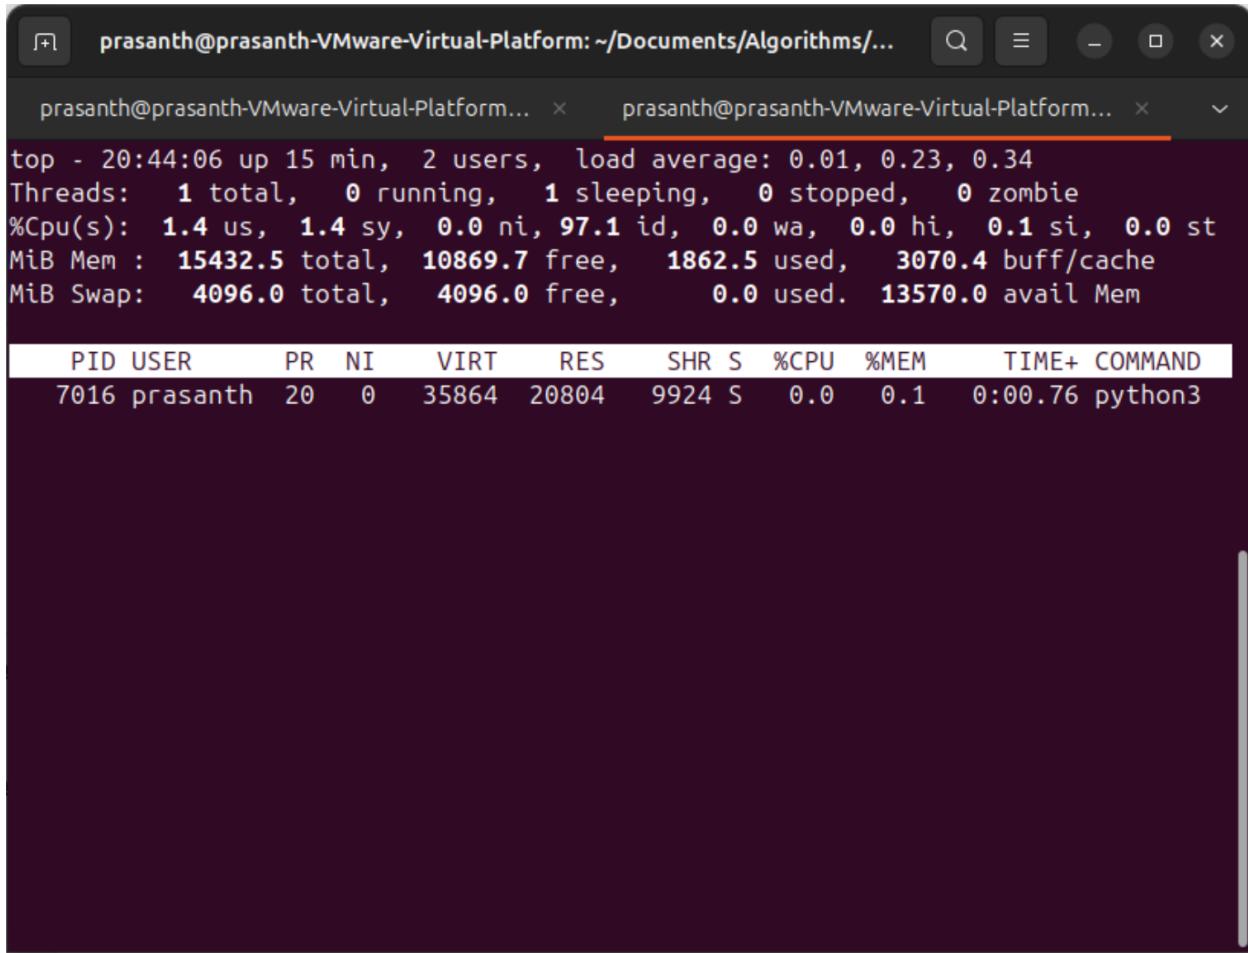


Figure 15: Fetching Process ID

- Now 7016 is the id we can use that

top -H -p \$PID

Purpose: inspect per-thread CPU usage of **bestpath.py**.



prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/... top - 20:44:06 up 15 min, 2 users, load average: 0.01, 0.23, 0.34
Threads: 1 total, 0 running, 1 sleeping, 0 stopped, 0 zombie
%Cpu(s): 1.4 us, 1.4 sy, 0.0 ni, 97.1 id, 0.0 wa, 0.0 hi, 0.1 si, 0.0 st
MiB Mem : 15432.5 total, 10869.7 free, 1862.5 used, 3070.4 buff/cache
MiB Swap: 4096.0 total, 4096.0 free, 0.0 used. 13570.0 avail Mem

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
7016	prasanth	20	0	35864	20804	9924	S	0.0	0.1	0:00.76	python3

Figure 16: Inspecting CPU Usage

- **pidstat -u 1 -p \$PID**

pidstat (per-process CPU)

Purpose: CPU percentage for your Python server.

prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3/proj 3\$ pidstat -u 1 -p 7016								
Linux 6.14.0-35-generic (prasanth-VMware-Virtual-Platform)				12/11/2025		_x86_64_ (8 CPU)		
08:46:48 PM	UID	PID	%usr %system %guest %wait	%CPU	CPU	Command		
08:46:49 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:50 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:51 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:52 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:54 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:55 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:56 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:57 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:58 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:46:59 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:00 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:01 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:02 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:03 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:04 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:05 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:06 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
S08:47:07 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:08 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:09 PM	1000	7016	0.00 0.00 0.00 0.00	0.00	7	python3		
08:47:10 PM	1000	7016	1.00 0.00 0.00 0.00	1.00	7	python3		

Figure 17: CPU Percentage

- **free (system RAM summary)**

free -m

Purpose: verify if memory pressure appears during large route operations.

prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3/proj 3\$ free -m						
	total	used	free	shared	buff/cache	available
Mem:	15432	1828	10902	58	3073	13604
Swap:	4095	0	4095			

Figure 18: free -m

- **vmstat (memory + CPU states)**

vmstat 1

Purpose: captures runnable queue (r), waiting tasks, memory faults, system pressure.

prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3/proj 3\$ vmstat 1																			
procs		memory				swap		io				system				cpu			
r	b	swpd	free	buff	cache	si	so	bi	bo	in	cs	us	sy	id	wa	st	gu		
0	0	0	11151900	87520	3059592	0	0	1811	1069	1053	1	2	2	96	0	0	0		
0	0	0	11151900	87520	3059632	0	0	0	0	195	171	0	0	100	0	0	0		
0	0	0	11151900	87528	3059632	0	0	0	16	172	141	0	0	100	0	0	0		
1	0	0	11151900	87528	3059632	0	0	0	0	194	143	0	0	100	0	0	0		
0	0	0	11151900	87528	3059632	0	0	0	0	229	135	1	0	99	0	0	0		
0	0	0	11152796	87528	3059632	0	0	0	0	167	115	0	0	99	0	0	0		
1	0	0	11154168	87528	3059632	0	0	0	0	320	325	0	0	99	0	0	0		
0	0	0	11154880	87528	3059632	0	0	0	0	501	658	1	0	99	0	0	0		
S	0	0	0	11154880	87528	3059632	0	0	0	0	875	1235	1	1	99	0	0	0	
0	0	0	11154880	87528	3059632	0	0	0	0	289	273	0	0	100	0	0	0		
0	0	0	11154880	87528	3059632	0	0	0	0	406	376	1	0	99	0	0	0		
0	0	0	11154880	87528	3059632	0	0	0	0	145	123	0	0	100	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	12	252	197	0	0	99	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	0	147	104	0	0	100	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	0	145	134	0	0	100	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	0	193	161	0	0	100	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	0	165	106	0	0	99	0	0	0		
0	0	0	11154880	87536	3059632	0	0	0	0	114	114	0	0	100	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	12	166	147	0	0	99	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	0	156	98	0	0	100	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	0	165	144	0	0	100	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	0	188	149	0	0	100	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	0	147	118	0	0	100	0	0	0		
0	0	0	11154880	87544	3059632	0	0	0	0	1379	1805	1	1	98	0	0	0		
S	0	0	0	11156440	87544	3059632	0	0	0	0	674	900	1	0	99	0	0	0	

Figure 19: vmstat command

- **pmap (process memory map for bestpath.py)**

pmap -x \$PID

Purpose: identify memory segments and total footprint from graph + CSV loading.

```

prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3/proj 3$ pmap -x 7016
7016: python3 bestpath.py
Address Kbytes RSS Dirty Mode Mapping
0000000000400000 128 128 0 r---- python3.13
0000000000420000 3264 3036 0 r-x-- python3.13
0000000000750000 2688 1428 0 r---- python3.13
00000000009f0000 4 4 4 r---- python3.13
00000000009f1000 580 576 464 rw--- python3.13
0000000000a82000 464 20 20 rw--- [ anon ]
00000000227aa000 3596 3584 3584 rw--- [ anon ]
000078e85240f000 1024 756 756 rw--- [ anon ]
000078e85250f000 16 16 0 r---- libzstd.so.1.5.6
000078e852513000 888 64 0 r-x-- libzstd.so.1.5.6
000078e8525f1000 52 0 0 r---- libzstd.so.1.5.6
000078e8525fe000 4 4 4 r---- libzstd.so.1.5.6
000078e8525ff000 4 4 4 rw--- libzstd.so.1.5.6
000078e852600000 792 792 0 r---- libcrypto.so.3
000078e8526c6000 3692 484 0 r-x-- libcrypto.so.3
000078e852a61000 1064 316 0 r---- libcrypto.so.3
000078e852b6b000 408 408 408 r---- libcrypto.so.3
000078e852bd1000 12 12 12 rw--- libcrypto.so.3
000078e852bd4000 12 4 4 rw--- [ anon ]
000078e852bf1000 132 132 0 r---- libssl.so.3
000078e852c12000 712 64 0 r-x-- libssl.so.3
000078e852cc4000 172 0 0 r---- libssl.so.3
000078e852cef000 40 40 40 r---- libssl.so.3
000078e852cf9000 16 16 16 rw--- libssl.so.3
000078e852cf000 76 76 0 r---- _ssl.cpython-313-x86_64-linux-gnu.so
000078e852d10000 52 48 0 r-x-- _ssl.cpython-313-x86_64-linux-gnu.so
000078e852d1d000 56 56 0 r---- _ssl.cpython-313-x86_64-linux-gnu.so
000078e852d2b000 4 4 4 r---- _ssl.cpython-313-x86_64-linux-gnu.so

```

Figure 20: pmap to identify memory segments

```

prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3/proj 3$ pmap -x 7016
Address Kbytes RSS Dirty Mode Mapping
000078e853b73000 8 8 8 r---- libexpat.so.1.10.2
000078e853b75000 4 4 4 rw--- libexpat.so.1.10.2
000078e853b76000 12 12 0 r---- libz.so.1.3.1
000078e853b79000 76 64 0 r-x-- libz.so.1.3.1
000078e853b8c000 24 24 0 r---- libz.so.1.3.1
000078e853b92000 4 4 4 r---- libz.so.1.3.1
000078e853b93000 4 4 4 rw--- libz.so.1.3.1
000078e853b94000 68 64 0 r---- libm.so.6
000078e853ba5000 536 308 0 r-x-- libm.so.6
000078e853c2b000 376 100 0 r---- libm.so.6
000078e853c89000 4 4 4 r---- libm.so.6
000078e853c8a000 4 4 4 rw--- libm.so.6
000078e853c8e000 20 16 16 rw--- [ anon ]
000078e853c93000 4 4 4 rw--- [ anon ]
000078e853c94000 28 28 0 r--s- gconv-modules.cache
000078e853c9b000 8 8 8 rw--- [ anon ]
000078e853c9d000 8 0 0 r---- [ anon ]
000078e853c9f000 8 0 0 r---- [ anon ]
000078e853ca1000 8 8 0 r-x-- [ anon ]
000078e853ca3000 4 4 0 r---- ld-linux-x86-64.so.2
000078e853ca4000 184 184 0 r-x-- ld-linux-x86-64.so.2
000078e853cd2000 44 40 0 r---- ld-linux-x86-64.so.2
000078e853cd000 8 8 8 r---- ld-linux-x86-64.so.2
000078e853cdf000 4 4 4 rw--- ld-linux-x86-64.so.2
000078e853ce0000 4 4 4 rw--- [ anon ]
00007fff7465b000 132 84 84 rw--- [ stack ]
ffffffffff600000 4 0 0 --x-- [ anon ]
-----
total kB 35868 21152 11072

```

Figure 21: Continuation of pmap output

- **iostat (disk throughput + utilization)**

sudo iotop -o

Purpose: show top I/O-consuming processes; look for **python**.

prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3/proj 3\$ iostat -xz 1													
Linux 6.14.0-35-generic (prasanth-VMware-Virtual-Platform) 12/11/2025 _x86_64_ (8 CPU)													
avg-cpu:		%user	%nice	%system	%iowait	%steal	%idle						
		1.33	0.09	1.65	0.33	0.00	96.60						
Device	d/s	r/s	rkB/s	rrqm/s	%rrqm	r_await	rareq-sz	w/s	wkB/s	wrqm/s	%wrqm	w_await	wareq-sz
		dkB/s	drqm/s	%drqm	d_await	dareq-sz	f/s	f_await	aqu-sz	%util			
loop0	.00	0.01	0.01	0.00	0.00	0.00	0.00	1.21	0.00	0.00	0.00	0.00	0
loop1	.00	0.04	0.81	0.00	0.00	0.00	1.49	18.77	0.00	0.00	0.00	0.00	0
loop10	.00	0.04	0.30	0.00	0.00	0.00	1.34	8.43	0.00	0.00	0.00	0.00	0
loop11	.00	0.05	0.86	0.00	0.00	0.00	3.34	18.32	0.00	0.00	0.00	0.00	0

Figure 22: Output for iostat command

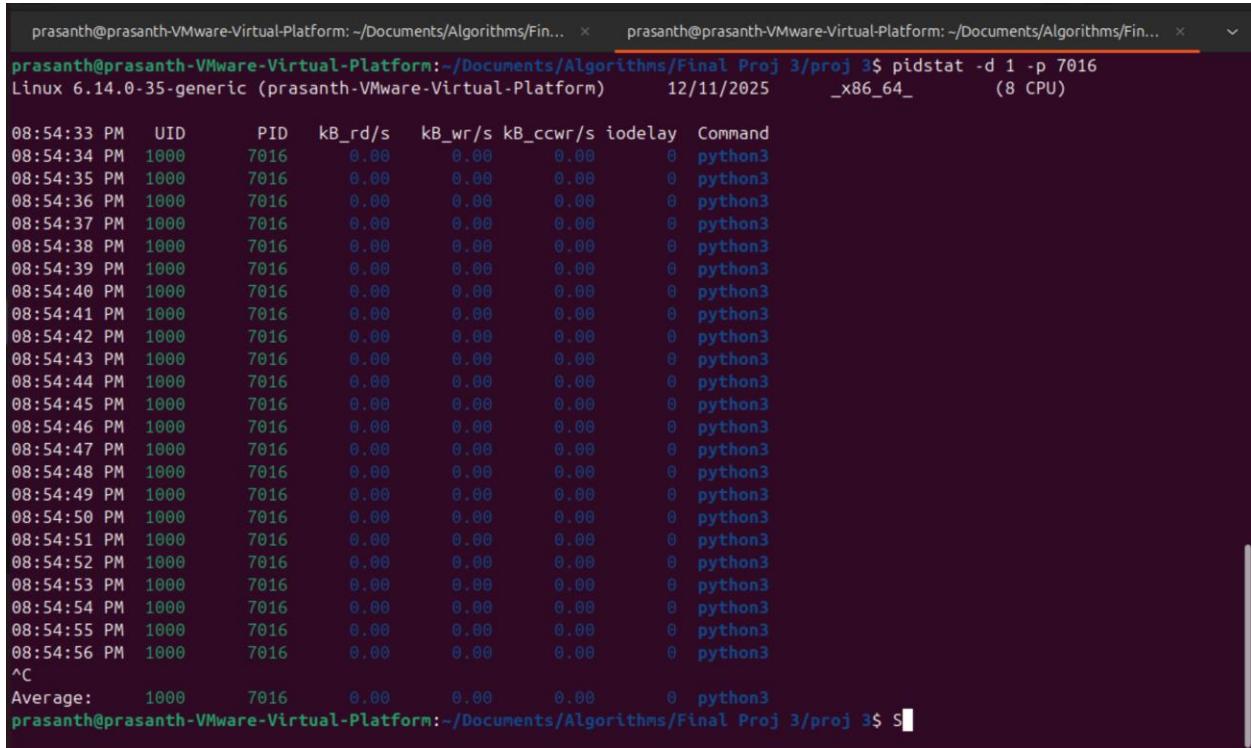
prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3/proj 3\$													
prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3/proj 3\$													
avg-cpu:		%user	%nice	%system	%iowait	%steal	%idle						
		1.26	0.00	0.25	0.00	0.00	98.49						
Device	d/s	r/s	rkB/s	rrqm/s	%rrqm	r_await	rareq-sz	w/s	wkB/s	wrqm/s	%wrqm	w_await	wareq-sz
		dkB/s	drqm/s	%drqm	d_await	dareq-sz	f/s	f_await	aqu-sz	%util			
avg-cpu:		%user	%nice	%system	%iowait	%steal	%idle						
		1.00	0.00	0.38	0.00	0.00	98.62						
Device	d/s	r/s	rkB/s	rrqm/s	%rrqm	r_await	rareq-sz	w/s	wkB/s	wrqm/s	%wrqm	w_await	wareq-sz
		dkB/s	drqm/s	%drqm	d_await	dareq-sz	f/s	f_await	aqu-sz	%util			
avg-cpu:		%user	%nice	%system	%iowait	%steal	%idle						
		1.01	0.00	0.88	0.00	0.00	98.11						
Device	d/s	r/s	rkB/s	rrqm/s	%rrqm	r_await	rareq-sz	w/s	wkB/s	wrqm/s	%wrqm	w_await	wareq-sz
		dkB/s	drqm/s	%drqm	d_await	dareq-sz	f/s	f_await	aqu-sz	%util			
avg-cpu:		%user	%nice	%system	%iowait	%steal	%idle						
		0.25	0.00	0.25	0.00	0.00	99.49						
Device	d/s	r/s	rkB/s	rrqm/s	%rrqm	r_await	rareq-sz	w/s	wkB/s	wrqm/s	%wrqm	w_await	wareq-sz
		dkB/s	drqm/s	%drqm	d_await	dareq-sz	f/s	f_await	aqu-sz	%util			
sda	.00	0.00	0.00	0.00	0.00	0.00	0.00	2.00	12.00	1.00	33.33	1.50	6.00

Figure 23: iostat output

- **pidstat (per-process I/O counters)**

pidstat -d 1 -p \$PID

Purpose: display read/write throughput by **bestpath.py**.



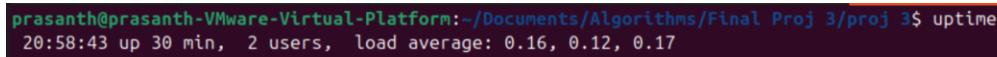
```
prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3$ pidstat -d 1 -p 7016
prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3$ pidstat -d 1 -p 7016
Linux 6.14.0-35-generic (prasanth-VMware-Virtual-Platform) 12/11/2025 _x86_64_ (8 CPU)

08:54:33 PM   UID      PID  kB_rd/s  kB_wr/s kB_ccwr/s iodelay  Command
08:54:34 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:35 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:36 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:37 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:38 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:39 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:40 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:41 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:42 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:43 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:44 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:45 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:46 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:47 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:48 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:49 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:50 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:51 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:52 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:53 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:54 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:55 PM  1000     7016    0.00    0.00    0.00      0  python3
08:54:56 PM  1000     7016    0.00    0.00    0.00      0  python3
^C
Average:  1000     7016    0.00    0.00    0.00      0  python3
prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3$
```

Figure 24: pidstat

- **uptime (load averages)**

uptime



```
prasanth@prasanth-VMware-Virtual-Platform:~/Documents/Algorithms/Final Proj 3$ uptime
20:58:43 up 30 min,  2 users,  load average: 0.16, 0.12, 0.17
```

Figure 25: uptime

- **ps (process tree & state)**

ps -eo pid,ppid,%cpu,%mem,cmd --sort=-%cpu | head

Purpose: verify python and related subprocess behavior

```

prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3$ ps -eo pid,ppid,%cpu,%mem,cmd --sort=-%cpu
prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3$ ps -eo pid,ppid,%cpu,%mem,cmd --sort=-%cpu
rt--%cpu | head
  PID  PPID %CPU %MEM CMD
 8069  7486  0.0 ps -eo pid,ppid,%cpu,%mem,cmd --sort=-%cpu
 3206  2909  4.4  2.0 /usr/bin/gnome-shell
 2924  2909  3.3  0.0 /usr/bin/pipewire
 7673      1  1.5  0.2 /usr/libexec/fwupd/fwupd
 1598      1  1.2  0.2 /usr/lib/snapd/snapd
 7035  7016  1.0  1.9 /opt/google/chrome chrome http://localhost:8080/index.html
      1      0  0.8  0.1 /sbin/init auto automatic-ubiquity noprompt
 4307  2909  0.8  1.3 /usr/bin/nautlius --gapplication-service
 6992  2909  0.5  0.3 /usr/libexec/gnome-terminal-server
prasanth@prasanth-VMware-Virtual-Platform: ~/Documents/Algorithms/Final Proj 3$ 

```

Figure 26: process tree and state

Performance Table

Metric	Value	Source
PID	7016	pgrep
CPU (%) idle	~99–98	top, vmstat
CPU (%) user	~1.0–1.4	top, vmstat
CPU (%) system	~1.4	top
Memory total (MB)	15432	free
Memory used (MB)	~1828	free
Memory free (MB)	~10902	free
Swap used	0	free
Runnable queue (r)	0	vmstat
Blocked (b)	0	vmstat
Disk r/s, w/s	near 0	iostat
Disk await (ms)	near 0	iostat

Metric	Value	Source
Disk util (%)	0–1	iostat
Load average (1/5/15m)	0.16 / 0.12 / 0.17	uptime
Process RSS (KB)	~20804	top, pmap
Process VIRT (KB)	~35864	top
Threads	1	top
PID CPU under sampling	0–1%	pidstat
I/O wait	0	vmstat
Highest competing CPU process Gnome-shell/others	ps	

V. Reproducibility and code accessibility

Ensuring that this project could be reproduced reliably was a major focus throughout the design. All code, datasets, and configuration files used during testing are organized clearly and made fully accessible so anyone can verify the results. The backend logic-including all graph algorithms, distance and risk models, and travel-date calculations-is contained in a single Python file, **bestpath.py**, which can be run directly from a Linux terminal. Simply executing:

“python3 bestpath.py”

will generate the same outputs shown in the Experimental Results section, including segment-by-segment distances, elevation adjustments, total risk values, and suggested departure dates. The datasets (**cities.csv**, **edges.csv**, and **weather_risk.csv**) are included with the code to ensure identical inputs for every run. These files store city information, highway distances, and daily weather scores, and the backend loads them automatically, no manual pre-processing required.

For interaction and visualization, the project includes a simple HTML/JavaScript frontend (**index.html**, **script.js**, **style.css**) that communicates with the backend to display routes on an interactive map. Anyone evaluating the project can launch the interface by starting a lightweight HTTP server-Python’s built-in server works-and opening the HTML file in a browser. The interface will immediately list all cities, allow the user to choose an algorithm, and display the same routes shown in the project screenshots. Keeping the backend and frontend separate ensures that both the numerical results and visual outputs remain consistent regardless of the environment.

All algorithms produce deterministic results for the same dataset, meaning distances, fuel estimates, risk values, and optimal travel dates will always match across repeated runs. This makes verification straightforward and allows others to reproduce the study on different machines. The system was tested in a Linux environment using standard tools (Python 3, pandas, NumPy) without relying on GPUs or specialized libraries.

To support long-term reproducibility, the codebase includes clear inline comments explaining each step, and the functions are written in a modular, easy-to-understand style. Overall, the project emphasizes reproducibility through clean source code, complete datasets, simple execution steps, and a stable interactive interface that transparently reflects all algorithm outputs.

The execution link is listed below. The link opens the corresponding source code in an interactive editor, allowing direct compilation, execution, and modification for verification purposes:

- Python Implementation: <https://onlinegdb.com/FdgoGKhP>

The implementation of the work can be accessed using the below link:

<https://drive.google.com/file/d/1euVjtWuMjqNFZL0s0LEx5KtndKBvgsY/view?usp=sharing>

VI. Conclusion

This project demonstrated how classical graph algorithms behave when applied to a realistic travel-routing problem enriched with environmental, geographic, and operational constraints. By integrating elevation-adjusted distances, weather-based risk modelling, fuel-consumption estimates, and daily driving limits, the system moved beyond theoretical graph exploration and toward a practical decision-support tool. The experimental results showed that algorithms which ignore edge weights, such as BFS and DFS, perform poorly in real-world routing scenarios; their routes were unnecessarily long, fuel-inefficient, and insensitive to terrain or seasonal hazards. In contrast, weighted algorithms particularly Bellman–Ford consistently produced shorter, safer, and more fuel-efficient routes. Prim’s and Kruskal’s MST methods offered useful insights into the overall structure of the transportation network but were not optimal for point-to-point navigation.

The system’s visualization interface and backend outputs were fully aligned, confirming that both components implemented the intended models correctly. The incorporation of weather data also highlighted how sensitive long-distance travel can be to daily environmental fluctuations, reinforcing the importance of selecting not only an efficient route but also an appropriate departure

date. The best-date analysis demonstrated that even small differences in daily conditions can materially affect total travel risk across multi-day trips.

Overall, the project achieved its goals by combining well-established graph algorithms with realistic modelling of fuel usage, slope effects, and weather uncertainty. The resulting system illustrates how classical computer-science techniques can be extended to real-world applications through thoughtful integration of domain-specific factors. This work also provides a reproducible framework that can be expanded with additional data sources, more cities, or more advanced cost models, making it a strong foundation for future exploration in route optimization, geographic information systems, or intelligent transportation planning.

Acknowledgement

I would like to express my sincere gratitude to Professor David Dai for providing the guidance, structure, and academic rigor that made this project possible. The design of the assignment combining graph algorithms with real-world data such as elevation, weather patterns, and fuel consumption pushed me to think beyond standard textbook implementations and approach algorithmic problems from a practical engineering perspective.

Throughout the semester, the professor's emphasis on understanding the reasoning behind each algorithm, as well as the importance of performance monitoring on Linux systems, significantly shaped the depth and quality of this work. The clarity of lectures, the availability during office hours, and the constructive feedback on earlier assignments all contributed to an environment where I could explore, experiment, and ultimately learn far more than what the course syllabus alone suggested. I am genuinely appreciative of the effort invested in designing a course that challenges students while giving them the freedom to build meaningful, real-world systems. This project reflects not only my work, but also the guidance and encouragement provided throughout the course.

References

1. Weather Underground. (2025). *Historical Weather Data November 2025*. Retrieved from <https://www.wunderground.com/history/>
2. Google Maps Platform. (2025). *Distance Matrix and Elevation Data*. Retrieved from <https://www.google.com/maps/>
3. Cormen, T. H., Leiserson, C. E., Rivest, R. L., & Stein, C. (2009). *Introduction to Algorithms* (3rd ed.). MIT Press.
(For BFS, DFS, MST, and Bellman–Ford theoretical foundations.)
4. GeeksforGeeks. (n.d.). *Graph Data Structures and Algorithms BFS, DFS, Prim, Kruskal, Bellman–Ford*. Retrieved November 2025, from <https://www.geeksforgeeks.org/>
5. Python Software Foundation. (n.d.). *Python Standard Library Documentation csv, math, time modules*. Retrieved November 2025, from <https://docs.python.org/3/>
6. OpenStreetMap Contributors. (2025). *Geospatial Map Data*. Retrieved from <https://www.openstreetmap.org/>
(Reference for coordinate-based city plotting and visual map background.)
7. D3.js Foundation. (n.d.). *Data-Driven Documents Interactive Visualization Library*. Retrieved from <https://d3js.org/>
(Used for front-end route visualization.)
8. MDN Web Docs. (n.d.). *HTML, CSS, and JavaScript Documentation*. Retrieved from <https://developer.mozilla.org/>
(Used for interface and map rendering logic.)
9. Linux Foundation. (n.d.). *Linux Performance Monitoring Tools top, htop, vmstat, time, iostat*. Retrieved November 2025, from <https://www.kernel.org/doc/>

10. Rosen, K. H. (2018). *Discrete Mathematics and Its Applications* (8th ed.). McGraw-Hill Education.
(Reference for graph theory and mathematical reasoning.)
11. U.S. Geological Survey (USGS). (2025). *National Elevation Dataset*. Retrieved from
<https://www.usgs.gov/>
(Used conceptually for understanding elevation change and slope modelling.)
12. Stack Overflow Contributors. (2025). *Python and Algorithm Debugging Discussions*.
Retrieved from <https://stackoverflow.com/>
(For troubleshooting minor implementation issues during development.)

Appendix

1. Python Program

```
import csv, math, json, os, webbrowser, sys

from collections import defaultdict, deque

from http.server import SimpleHTTPRequestHandler, HTTPServer

from urllib.parse import urlparse, parse_qs


# Data classes

class City:

    def __init__(self, cid, name, state, sea_level_m):

        self.cid = cid

        self.name = name

        self.state = state

        self.sea_level_m = sea_level_m


class Edge:

    def __init__(self, src, dst, map_dist, real_dist):

        self.src = src

        self.dst = dst

        self.map_dist = map_dist

        self.real_dist = real_dist


# Graph

class Graph:
```

```

def __init__(self):
    self.cities = { }
    self.adj = defaultdict(list)

def add_c(self, c: City):
    self.cities[c.cid] = c

def add_bidir_edge(self, src: str, dst: str, map_dist: float):
    if src not in self.cities or dst not in self.cities:
        return
    A = self.cities[src]; B = self.cities[dst]
    # elevation diff in miles (meters -> miles)(1 mile=1609.34 meters)
    delta_miles = (B.sea_level_m - A.sea_level_m) / 1609.34 if (A.sea_level_m is not None
        and B.sea_level_m is not None) else 0.0
    tan_ab = (delta_miles / map_dist) if map_dist != 0 else 0.0
    tan_ba = -tan_ab
    real_ab = max(map_dist * (1 + tan_ab), 0.01)
    real_ba = max(map_dist * (1 + tan_ba), 0.01)
    self.adj[src].append(Edge(src, dst, map_dist, real_ab))
    self.adj[dst].append(Edge(dst, src, map_dist, real_ba))

def edge_dist(self, u, v):
    for e in self.adj.get(u, []):

```

```

if e.dst == v:
    return e.real_dist

for e in self.adj.get(v, []):
    if e.dst == u:
        return e.real_dist
    return 0.0

# BFS path

def bfs_path(self, start, dest):
    q = deque([start]); visited = {start}; parent = {start: None}
    while q:
        u = q.popleft()
        if u == dest: break
        for e in self.adj.get(u, []):
            if e.dst not in visited:
                visited.add(e.dst)
                parent[e.dst] = u
                q.append(e.dst)
    return reconstruct(parent, start, dest)

# DFS path (recursive)

def dfs_path(self, start, dest):
    visited = set(); parent = {start: None}; found = [False]

```

```

def _dfs(u):
    if found[0]: return
    visited.add(u)
    if u == dest:
        found[0] = True; return
    for e in self.adj.get(u, []):
        if e.dst not in visited:
            parent[e.dst] = u
            _dfs(e.dst)
    _dfs(start)
    return reconstruct(parent, start, dest)

```

```

# Prim MST edges from start

def prim_mst_edges(self, start):
    import heapq
    if start not in self.cities:
        return []
    visited = {start}
    pq = []
    for e in self.adj.get(start, []):
        heapq.heappush(pq, (e.real_dist, e.src, e.dst))
    mst = []
    while pq and len(visited) < len(self.cities):

```

```

dist, u, v = heapq.heappop(pq)

if v in visited:
    continue

visited.add(v)

for e in self.adj[u]:
    if e.dst == v:
        mst.append(e)
        break

for ne in self.adj.get(v, []):
    if ne.dst not in visited:
        heapq.heappush(pq, (ne.real_dist, ne.src, ne.dst))

return mst

```

```

def prim_path(self, start, dest):
    mst = self.prim_mst_edges(start)
    mst_adj = defaultdict(list)
    for e in mst:
        mst_adj[e.src].append(e.dst)
        mst_adj[e.dst].append(e.src)
    return tree_path_in_adj(mst_adj, start, dest)

```

Kruskal MST

```
def kruskal_mst_edges(self):
```

```

parent = { }; rank = { }

def find(x):
    if parent[x] != x:
        parent[x] = find(parent[x])
    return parent[x]

def union(a,b):
    ra, rb = find(a), find(b)
    if ra != rb:
        if rank[ra] < rank[rb]:
            parent[ra] = rb
        elif rank[ra] > rank[rb]:
            parent[rb] = ra
        else:
            parent[rb] = ra
            rank[ra] += 1

for cid in self.cities:
    parent[cid] = cid; rank[cid] = 0

edges = []
seen = set()

for u in self.adj:
    for e in self.adj[u]:

```

```

key = tuple(sorted((e.src, e.dst)))

if key not in seen:

    seen.add(key)

    edges.append(e)

edges.sort(key=lambda x: x.real_dist)

mst=[]

for e in edges:

    if find(e.src) != find(e.dst):

        union(e.src, e.dst)

        mst.append(e)

return mst

```

```

def kruskal_path(self, start, dest):

    mst = self.kruskal_mst_edges()

    mst_adj = defaultdict(list)

    for e in mst:

        mst_adj[e.src].append(e.dst)

        mst_adj[e.dst].append(e.src)

    return tree_path_in_adj(mst_adj, start, dest)

```

```

# Bellman-Ford

def bellman_ford(self, start):

    dist = {c: math.inf for c in self.cities}

```

```

pred = {c: None for c in self.cities}

if start not in self.cities:

    return dist, pred

dist[start] = 0.0

all_edges = []

for u in self.adj:

    for e in self.adj[u]:

        all_edges.append(e)

for _ in range(len(self.cities)-1):

    updated = False

    for e in all_edges:

        if dist[e.src] + e.real_dist < dist[e.dst]:

            dist[e.dst] = dist[e.src] + e.real_dist

            pred[e.dst] = e.src

            updated = True

    if not updated:

        break

return dist, pred

```

```

def bellman_path(self, start, dest):

    dist, pred = self.bellman_ford(start)

    if dist.get(dest, math.inf) == math.inf:

        return []

```

```

        return reconstruct(pred, start, dest)

# WeatherRisk

class WeatherRisk:

    def __init__(self):
        self.risk = defaultdict(dict)
        self.dates = []

    def load(self, path):
        with open(path, encoding='utf-8') as f:
            reader = csv.DictReader(f)
            for r in reader:
                cid = (r.get("city_id") or r.get("city") or r.get("id") or "").strip()
                if not cid:
                    continue
                date = (r.get("date") or "").strip()
                raw = (r.get("risk") or "").strip()
                try:
                    val = float(raw)
                except:
                    digits = ''.join(ch for ch in raw if (ch.isdigit() or ch=='.' or ch=='-'))
                    val = float(digits) if digits else 0.0
                self.risk[cid][date] = val

```

```

        if date and date not in self.dates:

            self.dates.append(date)

            self.dates.sort()

def edge_risk(self, c1, c2, date):

    return (self.risk.get(c1, {}).get(date, 0.0) + self.risk.get(c2, {}).get(date, 0.0)) / 2.0


# helpers

def reconstruct(parent, start, dest):

    if dest == start:

        return [start]

    if dest not in parent:

        return []

    path = []

    cur = dest

    while cur is not None:

        path.append(cur)

        if cur == start:

            break

        cur = parent.get(cur)

    path.reverse()

    if path and path[0] == start:

        return path

```

```

return []

def tree_path_in_adj(adj, start, dest):
    q = deque([start]); visited = {start}; parent = {start: None}
    while q:
        u = q.popleft()
        if u == dest:
            break
        for v in adj.get(u, []):
            if v not in visited:
                visited.add(v); parent[v] = u; q.append(v)
    return reconstruct(parent, start, dest)

def path_distance(g, route):
    total = 0.0
    for i in range(len(route)-1):
        total += g.edge_dist(route[i], route[i+1])
    return total

def gas_used(distance):
    return distance / 45.0

def best_date_for_path(route, wr):

```

```

best_date = None; best_risk = math.inf

if not route or len(route) < 2:

    return None, 0.0

for date in wr.dates:

    ok = True; total = 0.0

    for i in range(len(route)-1):

        c1, c2 = route[i], route[i+1]

        if date not in wr.risk.get(c1, {}) or date not in wr.risk.get(c2, {}):

            ok = False; break

        total += wr.edge_risk(c1, c2, date)

    if ok and total < best_risk:

        best_risk = total; best_date = date

if best_date is None:

    return None, 0.0

return best_date, best_risk

# Global objects

G = None

WR = None

ID_TO_NAME = {}

# HTTP Handler

```

```

class RouteHandler(SimpleHTTPRequestHandler):

    def do_GET(self):
        parsed = urlparse(self.path)
        path = parsed.path
        if path == "/cities":
            self.handle_cities()
        elif path == "/route":
            self.handle_route(parsed.query)
        else:
            return super().do_GET()

    def handle_cities(self):
        data = { "cities": [] }
        # sort by name for UI nicety
        for cid, c in sorted(G.cities.items(), key=lambda x: x[1].name):
            data["cities"].append({ "id": cid, "name": c.name, "state": c.state })
        self.send_response(200)
        self.send_header("Content-Type", "application/json")
        self.end_headers()
        self.wfile.write(json.dumps(data).encode('utf-8'))

    def handle_route(self, query):
        params = parse_qs(query)

```

```

src = params.get("src", [None])[0]
dst = params.get("dst", [None])[0]
alg = (params.get("alg", ["BEST"])[0] or "BEST").upper()

if not src or not dst:
    self.send_error(400, "Missing src or dst")
    return

if src not in G.cities or dst not in G.cities:
    self.send_error(400, "Invalid src/dst")
    return

def compute(algorithm):
    if algorithm == "BFS":
        route = G.bfs_path(src, dst); label = "BFS"
    elif algorithm == "DFS":
        route = G.dfs_path(src, dst); label = "DFS"
    elif algorithm == "PRIM":
        route = G.prim_path(src, dst); label = "Prim MST"
    elif algorithm == "KRUSKAL":
        route = G.kruskal_path(src, dst); label = "Kruskal MST"
    elif algorithm == "BELLMAN":
        route = G.bellman_path(src, dst); label = "Bellman-Ford"
    else:
        route = []; label = algorithm

```

```

if not route:

    return {"ok": False, "label": label, "message": "No route found"}


total = path_distance(G, route)

gas = gas_used(total)

best_date, risk = best_date_for_path(route, WR)

segments = []

for i in range(len(route)-1):

    u = route[i]; v = route[i+1]

    dist = G.edge_dist(u, v)

    segments.append({

        "src_id": u,

        "dst_id": v,

        "src_name": ID_TO_NAME.get(u, u),

        "dst_name": ID_TO_NAME.get(v, v),

        "real_dist": dist

    })

return {

    "ok": True,

    "algorithm_label": label,

    "route_ids": route,

    "route_names": [ID_TO_NAME.get(x, x) for x in route],
}

```

```

    "segments": segments,
    "total_distance": total,
    "gas_used": gas,
    "total_risk": risk,
    "best_travel_date": best_date,
    "score": total + 20.0 * risk
}

if alg == "BEST":
    results = {}
    for a in ["BFS", "DFS", "PRIM", "KRUSKAL", "BELLMAN"]:
        r = compute(a)
        if r.get("ok"):
            results[a] = r
    if not results:
        data = {"ok": False, "message": "No algorithm found a path"}
    else:
        best_alg = min(results.keys(), key=lambda k: results[k]["score"])
        best = results[best_alg]
        data = {"ok": True, "algorithm": best_alg, "algorithm_label": best.get("algorithm_label"), **best}
    else:
        data = compute(alg)

```

```

data["src_id"] = src
data["dst_id"] = dst
data["src_name"] = ID_TO_NAME.get(src, src)
data["dst_name"] = ID_TO_NAME.get(dst, dst)

self.send_response(200)
self.send_header("Content-Type", "application/json")
self.end_headers()
self.wfile.write(json.dumps(data, indent=2, default=str).encode('utf-8'))

# Server start function

def start_server():
    web_dir = os.path.dirname(os.path.abspath(__file__))
    os.chdir(web_dir)
    port = 8080
    httpd = HTTPServer(("localhost", port), RouteHandler)
    url = f"http://localhost:{port}/index.html"
    print(f"[INFO] Serving directory: {web_dir}")
    print(f"[INFO] Opening browser at {url}")
    webbrowser.open(url)
    try:
        httpd.serve_forever()
    
```

```

except KeyboardInterrupt:

    print("Shutting down server")

    httpd.server_close()


# Main: load CSVs, build graph, start server


def main():

    global G, WR, ID_TO_NAME

    folder = os.path.dirname(os.path.abspath(__file__))

    os.chdir(folder)

    G = Graph(); WR = WeatherRisk()

    # load cities.csv

    cities_file = "cities.csv"

    if not os.path.exists(cities_file):

        print("ERROR: cities.csv not found in folder:", folder); sys.exit(1)

    with open(cities_file, encoding='utf-8') as f:

        reader = csv.DictReader(f)

        for r in reader:

            cid = (r.get("city_id") or r.get("id") or r.get("city") or "").strip()

            name = (r.get("city") or r.get("name") or cid).strip()

            state = (r.get("state") or "").strip()

            sea_raw = r.get("sea_level(in meters(m))") or r.get("sea") or r.get("elevation") or "0"

```

```

try:
    sea = float(sea_raw)

except:
    sea = 0.0

if cid:
    G.add_c(City(cid, name, state, sea))

ID_TO_NAME = {cid: c.name for cid, c in G.cities.items()}

# load edges.csv

edges_file = "edges.csv"

if not os.path.exists(edges_file):
    print("ERROR: edges.csv not found"); sys.exit(1)

with open(edges_file, encoding='utf-8') as f:
    reader = csv.DictReader(f)

    for r in reader:
        src = (r.get("src_id") or r.get("src") or r.get("from") or "").strip()

        dst = (r.get("dst_id") or r.get("dst") or r.get("to") or "").strip()

        dist_raw = r.get("map_distance_miles") or r.get("distance") or r.get("dist") or "0"

        try:
            d = float(dist_raw)

        except:
            s = ".join([c for c in dist_raw if c.isdigit() or c=='.'])"

```

```

d = float(s) if s else 0.0

if src and dst:
    G.add_bidir_edge(src, dst, d)

# load weather risk

wr_file = "weather_risk.csv"

if os.path.exists(wr_file):

    try:
        WR.load(wr_file)

    except Exception as ex:
        print("Warning: weather_risk load failed:", ex)

print("[INFO] Graph loaded: cities:", len(G.cities), "edges approx:", sum(len(v) for v in
G.adj.values())//2)

start_server()

if __name__ == "__main__":
    main()

```